

CeTTa Project Roadmap:

A Territory Map for MeTTa Runtime Design

Zar Goertzel

Oruži

Living document, started 2026-04-08

Abstract

This document describes the *territory* of programming language design decisions relevant to CeTTa, a direct C runtime for MeTTa. It is not a checklist of steps to complete, but a map of the landscape: regions to explore, peaks worth climbing, valleys to avoid, and clouds of possibility where the best path remains uncertain. We begin with an idealized vision of what CeTTa could become at its very best, then survey the terrain of term representation, evaluation semantics, search control, and storage. The journey so far is described as a continuous narrative, and future directions are presented as open questions rather than prescribed stages.

Contents

I	The Ideal	3
1	What Is the Very Best CeTTa Can Be?	3
1.1	The Vision	4
1.2	One Semantic Contract, Many Execution Lanes	4
1.3	What We Don't Know	5
1.4	What CeTTa Users Can Do	5
II	The Landscape	6
2	Term Representation and Identity	6
2.1	Why Representation Comes First	6
2.2	Hash-Consing: The Foundation	6
2.3	SymbolId: A Representational Cutover	7
2.4	The One-Universe Question	7
2.5	Canonical Storage ABI and Backend Path ABI	7
2.6	Clouds of Possibility	8
2.7	Reflective Evaluation	8
3	Evaluation Semantics	9
3.1	HE's Principal Evaluation Stages	9
3.2	The Contextual Nature of "Value"	9
3.3	TCO: A Constraint, Not a Checkbox	10
3.4	Unification and Cyclic Terms	10
3.5	Clouds of Possibility	10

4	Search and Control	11
4.1	WAM Foundations	11
4.2	Tabling: The Ladder of Capabilities	11
4.3	Current Seams	12
4.4	Search Fairness and Completeness	12
4.5	Parallel and Concurrent Execution	13
4.6	Bounded Demand and Streaming	13
4.7	Clouds of Possibility	14
4.8	Result Ordering	14
5	Storage and Indexing	14
5.1	Discrimination Trees and Substitution Trees	14
5.2	One Universe, Many Spaces	15
5.3	Revision-Aware Caching	15
5.4	Data Integrity and Security	15
5.5	Multiset Semantics over Set Storage	15
5.6	Prefix-Based Space Mirroring	17
6	Formal Backing	18
6.1	The Layered Confidence Model	18
6.2	Observable Equivalence	18
6.3	Evaluator Path Theorems	19
6.4	Binding and Substitution Safety	19
6.5	Term and Storage Invariants	19
6.6	Tabling and Cache Theorems	19
III	The Journey	19
7	The Path We’ve Walked	19
7.1	Starting with HE Fidelity	19
7.2	Arena Allocation and SymbolId	20
7.3	SearchContext and ChoicePoint	20
7.4	TAIL_REENTER and TCO	20
7.5	Where We Stand	20
IV	The Horizon	20
8	Clouds of Possibility	20
8.1	Eval-Local Term Sharing	21
8.2	Full SLG Tabling	21
8.3	Dependent Types as Extension	21
8.4	-lang petta Mode	21
8.5	PLN/ATP as Optional Lanes	21
8.6	Rho-Calculus Lane	21

9	The Cloud of Forks	22
9.1	Representation Fork	22
9.2	Substitution Fork	22
9.3	Evaluator/Machine Fork	22
9.4	Tabling Fork	23
9.5	Ordering/Demand Fork	23
9.6	Backend/Storage Fork	23
9.7	Verification Fork	23
A	Research Annex	23
B	Source Corpus	24
B.1	Hyperon / MeTTa Semantics	24
B.2	Comparative Implementations	25
B.3	Search/Control/Tabling	25
B.4	Term Representation / Indexing	25
B.5	Evaluator/Machine Design	25
B.6	Optional Reasoning Lanes	26
B.7	Verification / Proof	26
C	The Dual-Target Architecture	26
C.1	Proposed Source Syntax $\rightarrow$ Elaborated $\rightarrow$ Proof + Runtime	26
C.2	Semantic Skeleton: CBPV First, MATT Later	27
C.3	Why This Matters	27
C.4	Layer 3 as the Runtime Bridge	28
C.5	Three Sharing Regimes	28
C.6	Seven-Mode Structure	28
C.7	Four Theorems	29
C.8	Search Parameterization	29
C.9	Rholang and Efficiency	30
D	HE MeTTa Specification Summary	30
D.1	Syntax (S-Expression Grammar)	30
D.2	Core Evaluation Functions	30
D.3	Matching Algorithm	31
D.4	Minimal MeTTa Instructions	31
D.5	Lean Formalization	31
D.6	MM2 Interface	31

Part I

The Ideal

1 What Is the Very Best CeTTa Can Be?

Before examining techniques, trade-offs, or current status, we ask the fundamental question: if we could design CeTTa perfectly, what would it look like? This is the *Way of the Dotts*—define the idealized endpoint first, then work backwards. The destination shapes the journey.

1.1 The Vision

At its best, CeTTa would be a **semantically authoritative, high-performance MeTTa runtime** with these properties:

HE Fidelity as Baseline. The Hyeron Experimental (HE) semantics is not an aspiration but the foundation. Every well-typed MeTTa program that runs correctly in HE should produce the same observable results in CeTTa. Where CeTTa extends beyond HE (and it may), those extensions are explicit and documented, never silent behavioral drift.

Three Cleanly Separated Concerns. The mature systems we admire—E Prover, XSB, SWI-Prolog, Vampire—all eventually separate:

- **Semantics:** what the language means, independent of how results are computed.
- **Search and control:** how the runtime explores the space of possible reductions, manages nondeterminism, and returns results.
- **Storage and indexing:** where terms live, how they are retrieved, and what backends support them.

These concerns interact, but they are not the same thing. A clean separation means you can improve indexing without touching semantics, or change search strategy without breaking storage invariants.

One Shared Term Universe. Every persistent term should have a stable identity owned by one canonical store. Spaces are views or membership layers over that universe, not separate physical heaps. This is what Schulz discovered after years of “a mess” with multiple mutable term banks in E Prover: the converged design uses one immutable term bank with multiple specialized indices.

Specialized Engines as Extensions, Not Core. PLN backward chaining, ATP integration, factor-graph demand propagation, MORK algebraic operations—these are valuable, but they are not the evaluator core. They plug in through explicit interfaces. The evaluator itself is one coherent substrate, not a patchwork of special cases.

Performance Competitive with PeTTa. On recursive search benchmarks that exercise the WAM’s strengths, CeTTa should be competitive with PeTTa while preserving HE semantics that PeTTa diverges from. This is an engineering challenge, not a compromise of correctness.

Formal Verification as Confidence Anchor. Key invariants should be backed by Lean theorems: term canonicity, evaluator path equivalence, binding safety. Not as checkboxes to tick, but as anchors that give us confidence when we’re uncertain.

1.2 One Semantic Contract, Many Execution Lanes

The best possible CeTTa is not “a faster C port of HE.” It is a **family of semantically equivalent execution engines** living under one contract. At minimum, we envision:

Reference evaluator: Small, clear, easy to compare with Meta-MeTTa / HE ideas. Used for semantic reference and differential testing.

Production evaluator: Preserves direct tail reentry for singleton deterministic positions; uses resumable frames only where semantically necessary; supports streaming, bounded demand, and multiple results.

Tabled evaluator: Adds SLG-style suspension/resumption for recursive and cyclic search. Can grow from revision-aware caching toward full variant/subsumptive tabling.

Compiled/derived machines: Bytecode interpreter, WAM-like compiled control, or generated specialized evaluators for hot paths. Highest performance ceiling, greatest implementation complexity.

The core obligation: **all lanes must agree on observable results where they overlap.** This is what makes it one system, not a patchwork.

1.3 What We Don’t Know

The vision above describes properties, not implementations. We have beliefs about how to achieve them, but we acknowledge uncertainty:

- **Eval-local term sharing:** Borrowing canonical identity into evaluation arenas could reduce memory pressure significantly. Or it could introduce lifetime complexity that costs more than it saves. The Conchon–Filliâtre weak-reference model [1] suggests a path, but we haven’t walked it yet.
- **The right tabling granularity:** Full SLG suspension/resumption/completion (XSB-style) is powerful but complex. Simpler revision-aware memoization may suffice for many workloads. We don’t know where the sweet spot lies.
- **Contextual normal-form:** The question “is this term a value?” depends on expected type, evaluation context, and space contents. The literature on dependent types (Martin-Löf, Coquand) treats this carefully; we’re still learning how it applies to MeTTa’s specific type discipline.
- **Demand propagation:** When a bounded consumer (like `select`) needs only one result, the producer should stop. How to thread this demand through nondeterministic evaluation is an open research question.

The roadmap that follows describes the *territory* where these questions live, not definitive answers.

1.4 What CeTTa Users Can Do

The vision above is architectural. But a runtime exists for users, not architects. At its best, CeTTa lets you:

- **Load very large knowledge bases**—millions of atoms across distributed backends (DAS-scale)—without reimplementing your application.
- **Run deep recursive backward chaining** without stack failure. Recursive PLN queries over large knowledge graphs should complete, not crash.

- **Get bounded-demand answers** without overproduction. If you need only the first result, the evaluator should not compute thousands.
- **Switch backends without changing semantics.** Move from in-memory testing to PathMap persistence to distributed MORK without altering your MeTTa code or expecting different answers.
- **Stream aggregations** without full materialization. Fold over results as they arrive, not after collecting them all.
- **Trust theorem-backed invariants** on critical seams. When a Lean theorem says binding projection is sound, you can rely on it.
- **Run concurrent queries** over immutable snapshots. Parallel exploration of a frozen space state should be safe and deterministic.

These are user powers, not implementation checkboxes. The endgame is not “we have tabling” but “you can query recursive knowledge without crashing.”

Part II

The Landscape

2 Term Representation and Identity

Every operation in a term-rewriting runtime—pattern matching, binding application, result collection, memoization—operates on terms. The representation of terms therefore determines the cost of everything built on top. This is not speculation; it is the consensus of Knuth, Ritchie, Warren, Schulz, and every systems researcher who has built a mature symbolic runtime.

2.1 Why Representation Comes First

If terms are not shared, every `bindings_clone`, every pattern match, every result-set entry pays a representation tax that no local optimization can eliminate. You can tune individual operations, but you cannot escape the fundamental cost of copying.

Schulz’s E Prover provides the clearest case study. It began with mutable shared terms and multiple term banks. This was, in Schulz’s words at IWIL 2025, “a mess.” The mature E design converged on *one immutable term bank* with garbage collection, cached rewriting, and multiple specialized retrieval indices. The lesson: the architecture you resist at the start is the one you’ll eventually need.

Zhu et al. (JuliaSymbolics, 2025) quantified this. They bolted hash-consing onto an existing dynamically-typed symbolic runtime and measured $3.2\times$ speedup and $2\times$ memory reduction on a real biological signaling model (1,122 ODEs, 24,388 reactions). XSB’s manual states it directly: “interned ground terms are stored once in a global area; when tables are declared `intern`, those ground terms need not be copied into and out of tables.”

2.2 Hash-Consing: The Foundation

Hash-consing is the technique of storing structurally identical terms exactly once, returning the same pointer for equal structures. The benefits compound: equality becomes pointer comparison,

copying becomes aliasing, and memory usage becomes proportional to structural diversity rather than usage count.

The canonical reference is Conchon and Filliâtre’s 2006 paper on type-safe, modular hash-consing in OCaml. Their key insight: strong references in a hash table prevent collection and create dangling pointers when referents die. The solution is weak-reference hash tables where the garbage collector sweeps dead entries.

For a C runtime without GC, the question becomes: how do we get the benefits of hash-consing without the lifetime bugs? Two approaches appear in the literature:

Persistent-only hash-consing. Terms in the persistent arena are hash-consed; evaluation-local terms are not. Simple, but pays the copy tax during evaluation.

Arena-disciplined hash-consing. Evaluation arenas can borrow canonical terms from the persistent universe. Terms created during evaluation are local to that arena; when the arena resets, they vanish. Terms that need to persist are “promoted” to the persistent arena (and hash-consed there) before the evaluation arena dies.

The second approach is more powerful but requires careful lifetime discipline. An `ArenaResetSafety` theorem—formalizing that published terms survive arena rewind—would give confidence before attempting it.

2.3 SymbolId: A Representational Cutover

Profiling in March 2026 showed ~40% of CPU time in symbol string operations: interning, checking atom type, comparing strings. The fix was not a local optimization but a *representational cutover*: `SymbolId (uint32_t)` for global spellings, strings only at I/O boundaries.

This mirrors Rust’s own `Symbol` (interned index) and first-order prover functor-code disciplines (Vampire, E). The result: 38% instruction reduction on typical workloads.

2.4 The One-Universe Question

The ideal described in Part I calls for one canonical term universe with spaces as views over it. Currently, `Space` still owns `Atom **atoms`, and the `TermUniverse/term_canon` seam is a storage-policy shim rather than the one true store.

An important nuance from PathMap investigation: because PathMap is a prefix-sharing trie, naively encoding space membership as a leading prefix can *reduce* cross-space sharing (two otherwise identical terms diverge at the first path component). Canonical term storage should be separated from space membership indexing.

2.5 Canonical Storage ABI and Backend Path ABI

The persistent term store should be byte-packed, not pointer-heavy. This follows the E Prover / Vampire model: perfectly shared immutable term DAGs with cached immutable metadata.

Four ABIs. Do not collapse these into one format:

1. **Logical identity ABI:** `AtomId` over shared `TermUniverse`
2. **Canonical store ABI:** byte-packed immutable term records

3. **Evaluation ABI:** decoded/tagged hot-path views
4. **Backend path ABI:** self-contained structural bytes for interop

Internal vs External Bytes. The internal canonical record uses child `AtomIds` (DAG sharing). The external path key uses recursive structural encoding (portable). Do not confuse these—they optimize different concerns.

Models to Follow.

- E/Vampire/XSB for semantic/storage architecture
- YAP/WAM for hot evaluator ergonomics
- Cap’n Proto/FlatBuffers for byte record discipline
- PathMap for backend path keys, not as sole canonical truth

Born-Canonical Construction (Phase 1f). The “born-canonical” principle: terms should be constructed in canonical form at the store level, not built transiently and then re-interned. This follows Conchon–Filliâtre smart constructors: sharing is enforced at construction boundaries, and you never hash a term whose children aren’t already interned. The 16-byte `CettaTermHdr` format (tag, subtag, arity_or_len, sym_or_head, aux32, hash32) is fixed; Phase 1f adds the constructor substrate that produces these records directly.

2.6 Clouds of Possibility

Persistent-only vs. eval-local sharing: The conservative path hash-conses only persistent terms; the aggressive path borrows canonical identity into evaluation. We don’t know which is better for CeTTa’s workloads. Positive witnesses would include: backward chaining benchmarks completing where they currently OOM, hashcons-hit counters rising, and memory usage approaching PeTTa’s levels on comparable workloads.

Explicit substitutions: Krishnaswami (citing Abadi et al.) proposes deferring `bindings_apply` by representing `(term, substitution)` closures instead of eagerly copying. This eliminates the most expensive per-call operation but requires the evaluator to handle closures throughout. A long-term direction, not a near-term change.

2.7 Reflective Evaluation

MeTTa can quote and evaluate arbitrary expressions, including expressions that describe its own evaluator. This creates what Smith [20] called a *reflective tower*: a conceptually infinite stack of meta-circular interpreters, each interpreted by the one above.

Amin and Rompf [21] showed that such towers can be *collapsed*—stage polymorphism allows interpreters to compose in a way that eliminates interpretive overhead. Their Pink and Purple languages demonstrate that arbitrarily deep self-interpretation can be optimized away.

For CeTTa, reflection raises practical questions:

- **Tabling and self-modification:** If a tabled query’s definition changes mid-evaluation (via `add-atom`), what happens to cached answers? This is not hypothetical—PLN rule learning can modify the space while reasoning proceeds.

- **Meta-circular invariants:** When MeTTa evaluates a MeTTa evaluator, which invariants (typing, binding safety) transfer to the inner evaluation?
- **Collapsing for performance:** Can we detect when nested `eval!` calls are semantically equivalent to direct evaluation and optimize accordingly?

We do not have answers. But ignoring reflection would be ignoring one of MeTTa’s distinctive features. The reflective tower is not a bug or an edge case; it is the language’s design center for AGI self-improvement scenarios.

3 Evaluation Semantics

The evaluator is the heart of the runtime. It takes an atom (or outcome) and produces a stream of reduced atoms. Everything else—REPL, file loading, FFI—feeds into or out of the evaluator.

3.1 HE’s Principal Evaluation Stages

The HE specification in `metta.md` describes evaluation through mutually recursive functions. The principal stages are:

metta: Entry point; handles special forms, dispatches to interpretation.

interpret_expression: Inspects types, checks function-type applicability, dispatches.

interpret_function: Typed head/argument evaluation.

interpret_args: Evaluates arguments according to their declared strictness.

interpret_tuple: Handles tuple construction.

metta_call: Grounded execution and equation lookup.

These are the actual HE names. Prose labels like “EvalAtom” or “InterpretExpression” are fine for discussion, but the spec names matter when checking fidelity.

The key insight: `if` should not be a hardcoded C handler. The function type `(: if (-> Bool Atom Atom $t))` declares that the condition (`Bool`) is strict and the branches (`Atom`) are lazy. `interpret_args` reads the type and evaluates accordingly. The `stdlib` equation `(= (if True $t $e) $t)` then works through the general path. Per-op handlers are a tempting shortcut that becomes architectural debt.

3.2 The Contextual Nature of “Value”

When is a term fully reduced? The naive answer—“when no equations match”—is incomplete. In HE, the answer depends on:

- The **expected type**: an expression may be returned unchanged if the expected type is `Atom`.
- **Evaluated markers**: an expression already marked evaluated is not re-evaluated.
- **Atom category**: variables, errors, and empty are immediately returned.
- **Grounded dispatch**: a grounded function may apply even though no equations match.

- **Space contents:** the same term may reduce in one space and not another.

The dependent types literature (Martin-Löf, Coquand, Pfenning) treats normalization carefully because type checking depends on it. The principled answer comes from *normalization by evaluation* (NbE): evaluate into a semantic domain, then read back to syntax [22]. MeTTa’s type discipline is softer, but the same contextual structure applies. A bare predicate `IsValue(t)` is probably the wrong abstraction; something like `NormalFormAt(space, expectedType, atom, bindings)` captures the dependencies more faithfully.

3.3 TCO: A Constraint, Not a Checkbox

Tail-call optimization for deterministic single-result forms (`let`, `chain`, `let*`, `if`, `case`, `function/return`) is not an optional performance feature. It is a semantic constraint: deeply recursive MeTTa programs must not blow the C stack.

The `TAIL_REENTER` macro in CeTTa’s `eval.c` implements this as a `goto tail_call` trampoline. It is cheap, correct, and *must be preserved* in any refactoring. The core-opt epoch’s central mistake was replacing `TAIL_REENTER` for `let/chain` with full frame allocation—a semantic regression that produced wrong answers, not just slowdowns.

3.4 Unification and Cyclic Terms

What counts as a valid binding? Logic runtimes differ materially here. SWI-Prolog supports rational trees (cyclic structures) and exposes `unify_with_occurs_check/2` for those who want to reject cycles. XSB takes a different approach. This is not a performance detail—it changes what programs mean.

For CeTTa, this remains an **open design question**:

- Does CeTTa reject cyclic bindings by default? HE appears to reject them (no rational tree support), but whether this is fundamental or could be relaxed remains open.
- Should rational trees ever be supported as an opt-in mode?
- Is occurs-check always semantic, or configurable per-lane?
- How do compiled/tabled lanes preserve unification policy consistently?

The answer affects what matchers must reject, what tables can store safely, and what alpha-equivalence means. Settling this is prerequisite to serious tabling and compiled-lane work.

One promising discipline, especially for future relation-compiled lanes, is to split unification into two interfaces: a *guard predicate* (“does a match exist?”) and a deterministic *binding extractor* (“what substitution does this chosen match induce?”). Positive example: a rho-style COMM lane can let the guard participate in search/saturation while delaying substitution construction until a rewrite actually fires. Negative example: one monolithic unification hook that entangles search branching, fixpoint existence, and binding materialization.

3.5 Clouds of Possibility

Frame-based evaluation for typed-apply: Typed ordinary application with multi-step argument evaluation benefits from heap-allocated frames. The question is: can we isolate this to typed-apply only, preserving `TAIL_REENTER` for everything else?

Direct vs. stack equivalence: If we have both recursive and frame-based evaluation paths, they must produce identical results on the overlapping subset. A Lean theorem (`DirectVsStackEquivalence`) would anchor this.

Demand propagation: When `select` or `once` bounds the consumer, the producer should stop producing. This requires either generator-style suspension or explicit demand threading. Neither is fully designed.

4 Search and Control

MeTTa is nondeterministic: a single expression can reduce to multiple outcomes. Search and control determine how the runtime explores this space and returns results.

4.1 WAM Foundations

The Warren Abstract Machine (WAM) remains the gold standard for efficient Prolog execution after four decades. Its core ideas:

Choicepoints: Snapshots of state before nondeterministic choices.

Trail: A log of bindings made since the last choicepoint.

Backtracking: Undo the trail, restore the choicepoint, try the next alternative.

Last-call optimization: If no choicepoint remains in the current frame, reuse the frame for the tail call.

These ideas transfer to any nondeterministic language, though MeTTa’s type-directed evaluation adds complexity the WAM doesn’t face.

A crucial WAM distinction: **environments** (AND-nodes) are call frames holding local variables and return addresses; **choicepoints** (OR-nodes) are snapshots taken before nondeterministic choices [3]. Environments grow and shrink with deterministic calls; choicepoints freeze the stack until backtracking. In CeTTa, `SearchContext/ChoicePoint` map to OR-nodes, while `NativeResumeStack/EvalFrame` serve as AND-nodes for streaming evaluation. Keeping this distinction clear prevents frame-management bugs that plagued early implementations.

4.2 Tabling: The Ladder of Capabilities

Tabling (memoization of goals and their answers) eliminates redundant computation and handles left-recursive predicates that would loop in naive evaluation. XSB is the reference implementation; SWI-Prolog provides a modern alternative with delimited continuations. Stryer-Prolog implements tabling via the Desouter et al. approach [6].

The tabling landscape forms a **ladder of capabilities**:

Level 1: Revision-aware caching. Store query results keyed by (`space`, `revision`, `pattern`). Invalidate when the space changes. Simple, but coarse-grained—any change invalidates all entries for that space.

Level 2: Variant tabling. Cache based on structural equivalence of query patterns (ignoring variable names). `foo(X, Y)` and `foo(A, B)` share results. More reuse than simple memoization.

Level 3: Suspension/resumption/completion. Consumers block on incomplete tables; producers resume when new answers arrive. When all producers finish, tables are marked complete. This is full SLG resolution—required for recursive queries that don’t terminate under naive evaluation.

Level 4: Subsumptive tabling. Reuse answers from more general queries for more specific instances. If `foo(X)` is cached, `foo(bar)` can reuse those results filtered by unification. Greater sharing, more complex implementation.

Level 5: Answer subsumption. Remove redundant cached answers when one subsumes another. Important for lattice-valued or aggregating queries (shortest path, best proof).

Level 6: Incremental tabling. Track dependencies between facts and cached queries. When a fact changes, invalidate only transitively-dependent entries. Essential for applications with frequent updates (dynamic knowledge graphs, belief refinement).

Level 7: Monotonic/shared/parallel tabling. Monotonic tabling restricts queries to return monotonically increasing sets (terminates earlier). Shared tables across threads enable parallel search. These are advanced features in mature systems.

CeTTa need not climb the full ladder immediately. For many workloads, Level 1–2 suffice. But for recursive PLN chaining over large knowledge bases, Level 3+ becomes essential. The question is: which workloads justify which levels?

4.3 Current Seams

CeTTa’s `SearchContext` and `ChoicePoint` structures in `search_machine.h` are real, small abstractions. They’re not full SLG, but they’re not nothing either. `TableStore` provides revision-aware memoization infrastructure.

The `OrderedOutcomeVisitor` and ordered-visitation seam handle streaming results in the current evaluator. This is the path where `collapse`, `fold`, and `select` operate.

The next refinement above these seams is a single explicit episode layer: a `SearchEpisode` owning a normalized `SearchRequest` and a `SearchLangContract`, dispatching through a `SearchMachineOps` vtable to the existing `SearchContext` for branch-local state, to `SpaceMatchBackendOps` for storage, and to `TableStore` (wrapped as the first `SearchTableOps` plugin) for memoization. The seam is read/search only in v1—lane families (MM2-step, ATP saturation, solver-oracle) belong to a later epoch. The implementation plan carries this as Phase 2 with explicit non-goals.

4.4 Search Fairness and Completeness

Depth-first search is simple and memory-efficient, but incomplete: it can loop forever on infinite branches, never reaching solutions on other paths. Breadth-first search is complete—it finds every solution—but uses memory exponential in depth. This is not an abstract concern; recursive PLN queries can easily produce infinite search spaces.

The literature offers several solutions:

Iterative deepening [23]: Run depth-limited DFS with increasing limits. Achieves BFS completeness with DFS memory. The overhead of re-expansion is $\sim b/(b-1)$ where b is branching factor—acceptable for most workloads.

Interleaving (miniKanren) [24]: Interleave exploration of disjuncts, guaranteeing no branch starves indefinitely. Complete even with infinite branches, as long as some delay occurs in recursive cycles.

Bounded-depth with tabling: Use cycle detection in the table to prevent infinite loops. SLG resolution (Level 3) handles this naturally.

CeTTa currently uses essentially depth-first evaluation with some ordered-streaming structure. This is not a bug but a known limitation. The tabling ladder provides the path to completeness for recursive queries. Whether to adopt interleaving (which changes observable result order) or rely entirely on tabling (which adds infrastructure) remains an open design question.

4.5 Parallel and Concurrent Execution

The scale story matters. A runtime that can only use one core on a 128-core machine is leaving performance on the table. Options for parallelism:

Snapshot-parallel query: Freeze the space as an immutable snapshot, then run multiple query threads with thread-local scratch. No locks needed for reads; writes are forbidden. Simple and effective for read-heavy workloads.

Shared tabling: XSB-style thread-shared tables enable parallel producers feeding shared consumers. Coordination overhead exists but can be amortized over large result sets.

Rho-calculus concurrency: Meredith’s reflective higher-order process calculus [25] provides formal process-algebraic semantics for concurrent rewriting. MeTTa’s quote/unquote naturally maps to rho’s $\textcircled{}$ and $\ast$ operators, offering a principled foundation for parallel execution without ad-hoc concurrency semantics.

DAS distribution: The Distributed Atomspace [9] shards knowledge across machines. For million-atom scales, this is not optional but necessary.

A fundamental question: is MeTTa evaluation deterministic? If so, parallel execution is semantic-preserving and result order can be relaxed. If not (as with **random** or concurrent space mutation), the semantics of parallel evaluation must be carefully defined.

4.6 Bounded Demand and Streaming

Many queries need only partial results:

- **select k** — return first k answers, stop producing
- **once** — return first answer only
- **fold** — aggregate results incrementally without full materialization
- Streaming consumers — process answers as they arrive

These are not mere conveniences. For large nondeterministic searches, bounded demand is the difference between “finishes in milliseconds” and “runs forever computing unused results.”

This is not only a usability issue. In fragments with multiset or bag-style pattern matching (rho-style COMM, certain PLN conjunctions), enumerating all compatible matches is combinatorial in

the worst case. For those lanes, **once**, **select k**, and related demand bounds are part of the tractability contract, not just presentation knobs.

The challenge: threading demand through an eager evaluator. Options include generator-style suspension (yield and resume), explicit demand tokens propagated through choicepoints, or penguins-style incremental retrieval. None is fully designed for CeTTa.

4.7 Clouds of Possibility

Full SLG vs. simpler caching: Implementing full suspension/resumption/completion is substantial work. For which workloads would it matter? Recursive PLN chaining? Large knowledge-base queries? We should identify concrete witnesses before committing.

Demand termination: Bounded consumers should stop unbounded producers. This is easy to state, hard to implement in an eager evaluator. Lazy streaming or explicit demand tokens are two approaches.

Parallel search: Snapshot-parallel query over immutable space snapshots with thread-local scratch enables parallelism without locks. But MeTTa’s semantics currently assume deterministic ordering in some places. Relaxing this requires semantic clarity first.

4.8 Result Ordering

Result ordering depends on the search strategy and engine selected. CeTTa does not promise a fixed presentation order by default—only multiset equality of results. A deterministic-order mode could be offered as an explicit option for reproducibility-sensitive workloads.

The HE spec is authoritative here; CeTTa follows whatever ordering guarantees HE makes (or doesn’t make). If HE specifies deterministic order for certain operations, CeTTa must match. If HE leaves order unspecified, CeTTa may relax it for parallelism or other optimizations.

5 Storage and Indexing

Where do terms live? How are they retrieved? These questions intersect with term representation but are not identical.

5.1 Discrimination Trees and Substitution Trees

First-order theorem provers index their clause sets for fast retrieval. Two structures dominate:

Discrimination trees (E Prover). Index by functor, then by argument structure. Fast for forward reasoning and rewriting. Good when you want all terms matching a pattern.

Substitution trees (Graf 1994). Combine discrimination-tree speed with abstraction-tree generality. Crucially, they produce bindings during traversal, not just matching terms. This is exactly what pattern matching in MeTTa needs.

CeTTa’s `subst_tree.c` explicitly says it produces partial bindings during the walk—Graf-style.

5.2 One Universe, Many Spaces

The ideal: one canonical term universe, with spaces as membership layers indexed separately. The current reality: `Space` owns its own `Atom **atoms`, and migration hasn’t begun.

PathMap offers an alternative storage backend: a prefix-compressed trie where common prefixes share nodes. It supports algebraic operations (union, intersection, difference) that native arrays don’t.

The Distributed Atomspace (DAS) [9] represents another backend direction: very large knowledge bases across machines, with a MeTTa-integrated query API and multiple persistence engines. For million-atom scales, DAS-style distributed backends become essential.

The question for all backends is **semantic parity**: given the same logical space contents, do native, PathMap, MORK, and DAS backends produce identical evaluation results? This is a formal obligation, not just a testing goal.

5.3 Revision-Aware Caching

When evidence is added or removed, which cached results invalidate? The `space_revision` counter (bumped on every logical content change) supports lazy invalidation: cache entries include a revision number; lookup misses if the space has moved on.

This is coarse-grained (any change invalidates all entries for that space). Fine-grained invalidation (tracking which terms changed and which caches depended on them) is an open research area.

5.4 Data Integrity and Security

For distributed backends, integrity becomes load-bearing infrastructure:

Revision vectors: Causal ordering of space mutations. When multiple nodes can write, version vectors or Lamport timestamps establish which update happened “before” another. Without this, “eventual consistency” is meaningless.

Immutable logs: Append-only logs of facts provide provenance and enable replay. DAS uses this pattern; so do blockchain-inspired knowledge graphs.

Backend parity testing: Same logical contents must produce identical results across backends. This is not just a test suite but a formal obligation—the `BackendSemanticParity` theorem.

Access control: Who can read or write which spaces? For AGI systems with self-modification, this is not a nice-to-have but a safety requirement.

This is not “security theater.” A distributed knowledge base without integrity guarantees is a distributed bug generator.

5.5 Multiset Semantics over Set Storage

MeTTa has *multiset* semantics: adding the same atom twice creates two copies, and queries return one result per copy. PathMap and MORK use *set* semantics by design—duplicates are automatically removed. This is intentional for MM2 deterministic execution but creates a semantic mismatch when PathMap backs ordinary MeTTa spaces.

Current workaround. CeTTa’s bridge maintains a parallel `atom_index_rows` HashMap mapping serialized atoms to row-ID vectors. On insert, each copy gets a unique row ID; queries consult the parallel index for multiplicity. This is fragile: algebraic operations (`pjoin`, `pmeet`, `psubtract`) do not propagate to the parallel index, so combining spaces can silently corrupt multiplicities.

Approach A: Embedded count (works today). Encode multiplicity in a CeTTa-owned counted key formed by

`atom_expr_bytes || count_expr_bytes`

where the count is still stored as a normal serialized number expression. This keeps the atom bytes at the very front of the key so CeTTa can prefix-scan on atom structure and then decode the trailing count. Using the ordinary pair expression (`$atom $count`) directly would *not* be prefix-friendly under current MORK expr serialization, because the outer arity byte appears before the atom payload. On insert, look up the existing count and re-insert with `count+1`. On query, decode the count and emit the result that many times.

- **Pro:** Works with mainline PathMap, no API changes.
- **Con:** Count rewrites still live in CeTTa’s query/storage layer, and algebra ops still don’t compose cleanly.

Approach B: PathMap<CountValue> (clean, needs additions). Use `PathMap<CountValue>` where `CountValue` is a newtype wrapping a count. This requires:

1. Proper semiring `Lattice` impl for `CountValue`. Current numeric impls in `ring.rs` are placeholders that return `Identity` (keeping self, ignoring other).
2. Query-level count multiplication for conjunctions: storage gets cleaner under this model, but query evaluation must still multiply counts across matched factors.

Note that `join_val_at` is a convenience, not a strict blocker: `PathMap` already supports in-place insert-or-mutate via `write_zipper_at_path` (`trie_map.rs:450`).

Semiring algebra for multisets. Multiset (bag) semantics are well-studied in database theory. The canonical reference is Green et al. [16] on *provenance semirings*, which parameterize query evaluation over an arbitrary commutative semiring $(\mathbb{K}, +, \times, 0, 1)$. For counting multiplicities, the natural choice is $(\mathbb{N}, +, \times, 0, 1)$:

- **Union:** $\text{count}_a + \text{count}_b$ (additive).
- **Join/Product:** $\text{count}_a \times \text{count}_b$ (Cartesian product multiplicity).
- **Projection:** sums multiplicities over projected-away witnesses (not mere preservation).
- **Intersection:** $\min(\text{count}_a, \text{count}_b)$ (lattice meet).
- **Difference:** $\max(0, \text{count}_a - \text{count}_b)$ (saturating subtraction).

See also Mumick et al. [17] on bag semantics in Datalog and Abiteboul–Hull–Vianu [18] for bag relational algebra.

Benefits of clean semiring approach.

- Multiset semantics compose with all `PathMap` algebra operations.
- No parallel index to maintain (eliminates `atom_index_rows`).
- Conjunction multiplicity “just works” (product of factor counts).
- Clean separation: MORK uses `PathMap<()>` (set semantics), MeTTa spaces use `PathMap<CountValue>` (multiset semantics).
- Formal backing: semiring semantics are well-studied and can be formalized in Lean against the existing `PathMapQuantale` hierarchy.

CeTTa-owned tunnel for Approach A. Even before Approach B is implemented, CeTTa can already collapse the current generic pathmap-backed lane to a single trie-owned truth:

- `new-space pathmap` remains fully CeTTa-owned.
- PathMap stores one counted key per distinct atom.
- Insert/remove-one rewrite the counted key rather than maintaining a parallel duplicate-row index.
- Query/conjunction logic becomes a local CeTTa concern: prefix scan, local rematch, count readout, and factor-count multiplication.
- This removes the need for `atom_index_rows` and `query_multi_with_factor_exprs` in that backend lane.

If CeTTa ever deliberately mixes counted keys with other key families in one shared PathMap, a discriminator byte can be reintroduced then; the simplest fresh-space prototype does not need one.

Formalization opportunity. The Mettapedia formalization could extend `PathMapQuantale` with a `CountValue` instance and prove that semiring laws distribute over the trie operations. This would provide mechanically-verified multiset semantics for PathMap-backed spaces.

5.6 Prefix-Based Space Mirroring

A powerful view/elaboration pattern emerges from MM2’s prefix representation of spaces. In MORK, a space `&kb` can be represented as a prefix family (`kb $atom`):

```
; MM2 representation: prefix-tagged atoms
(kb (edge a b))
(kb (edge b c))
(kb (type-of a Person))

; Equivalent to HE/CeTTa:
(add-atom &kb (edge a b))
(add-atom &kb (edge b c))
(add-atom &kb (type-of a Person))
```

This suggests a **mirroring view layer**:

The idea. Mirror a large MORK space (potentially millions of atoms) into a prefix-filtered view that CeTTa can query efficiently:

- (`mirror &cetta-view &mork-space (kb $_)`) — create a view containing only (`kb ...`) prefixed atoms
- The view tracks a subset; updates propagate lazily
- As long as the view is small enough, mirroring stays fast
- CeTTa queries the view; MORK owns the full space

Quantale semantics. This aligns with PathMap’s `prestrict` operation: restricting a trie to paths whose prefix appears in another trie. The view is literally `prestrict(full_space, prefix_selector)`.

Implementation sketch.

1. Define `PrefixView` as a wrapper holding a reference to a MORK space plus a prefix pattern
2. On query, translate to MORK prefix-scan + local rematch

3. On mutation, propagate to underlying space with prefix prepended
4. Track revision numbers for lazy invalidation

When this wins.

- Large knowledge bases with structured prefixes (ontologies, taxonomies)
- Cross-system bridges where only a slice is needed
- Incremental reasoning over a stable background theory

This is a **roadmap item**, not current implementation. Source-syntax elaboration and backend separation from Phase 1f must stabilize first. The prefix view builds on top as an optimization layer.

6 Formal Backing

Key invariants should be mechanically verified, not just believed. This section surveys what formal obligations exist, where Lean formalization fits, and how confidence layers build upon each other.

6.1 The Layered Confidence Model

Verification is not all-or-nothing. We envision three layers:

Layer 1: Executable conformance. Cross-implementation test suites (e.g., `logicmoo/metta-testsuite`), differential testing against HE/PeTTa, property-based tests for evaluator equivalence, randomized backend parity tests. This is the foundation—without passing tests, nothing else matters.

Layer 2: Mechanized theorems. Lean proofs of load-bearing invariants: observable equivalence, TCO preservation, binding safety, hash-cons invariants, cache coherence, backend parity. These are confidence anchors: when a theorem is proved, we know with certainty that the property holds. When a theorem has `sorry`, we know exactly where our confidence stops.

Layer 3: Proof-producing reasoning. Selected external reasoners (ATP systems, specialized engines) return proof objects that CeTTa can consume or reconstruct. MeTTaMath and Lean integration provide bridges from reasoning to verification. This is the far horizon—genuinely exciting for AGI self-auditing scenarios.

More broadly, **derivation tracking** is an open question: can a user ask “how was this result derived?” Even E Prover has `-p` mode for proof output. Options range from simple unification traces to full proof terms. Whether this is optional (enabled by flag) or always available, and how much overhead it costs, remain design questions worth exploring.

6.2 Observable Equivalence

Two evaluation paths are equivalent if they produce the same result multiset (order may differ). This is the fundamental correctness criterion: if we have both recursive and frame-based evaluators, they must be observably equivalent.

6.3 Evaluator Path Theorems

DirectVsStackEquivalence: The recursive evaluator and the frame-based evaluator produce identical results on their shared domain.

TCOPreservation: The `TAIL_REENTER` path does not change semantics, only stack usage.

6.4 Binding and Substitution Safety

VisibleBindingProjection: Projecting only body-visible bindings does not change the outcome.

BranchDeltaMerge: Merging branch-local binding deltas back into the outer environment is sound.

TypeDirectedArgumentOrder: Evaluating arguments in type-directed order (strict first, lazy deferred) matches HE semantics.

6.5 Term and Storage Invariants

HashConsInvariant: Structurally equal terms share identity.

ArenaResetSafety: Published terms survive eval-arena rewind.

SpaceRevisionCoherence: Revision bumps and cache keys agree with logical mutation boundaries.

BackendSemanticParity: Same logical space contents produce identical results across backends.

6.6 Tabling and Cache Theorems

SoundCaching: Cached answers are correct answers (no false positives).

CompleteAnswers: When a table is marked complete, no additional answers exist.

IncrementalInvalidationSoundness: Dependency-based invalidation removes all affected entries.

Part III

The Journey

7 The Path We've Walked

This section describes what CeTTa has done, written as if we envisioned it from the start. The narrative is continuous because the work is continuous; we didn't know everything at the beginning, but each step followed from the previous.

7.1 Starting with HE Fidelity

The earliest decision: CeTTa would be a *direct* HE *runtime*, shaped around the HE evaluation structure. Not an interpreter for a different language that happens to parse MeTTa. Not a re-implementation that diverges where convenient. The HE spec is the contract.

This meant structuring `eval.c` around the mutual recursion of `metta`, `interpret_expression`, `interpret_function`, and friends. It meant threading expected type throughout, even when that

was awkward. It meant implementing type-directed argument evaluation even though a simpler eager strategy would have been faster in some cases.

7.2 Arena Allocation and SymbolId

The first representation work: arena allocation for evaluation-local terms, avoiding malloc/free overhead. Then the `SymbolId` cutover, replacing string comparisons with integer comparisons. 38% instruction reduction. These were low-hanging fruit, but they set the pattern: representation changes pay off more than local optimizations.

7.3 SearchContext and ChoicePoint

As nondeterminism became more important, we introduced `SearchContext` and `ChoicePoint` as explicit structures. Small, real abstractions—not full SLG, but the beginning of it. `TableStore` followed, providing revision-aware memoization.

7.4 TAIL_REENTER and TCO

The `TAIL_REENTER` pattern emerged from necessity: deeply recursive programs were blowing the stack. A `goto tail_call` trampoline for deterministic single-result forms solved it. Now 15 call sites in `eval.c` use this pattern.

7.5 Where We Stand

The current position on the map:

- **Term representation:** Arena allocation, `SymbolId`, persistent hashcons—all working. Eval-local hashcons not attempted (needs arena-reset-safety confidence first).
- **Evaluation:** Type-directed, `TAIL_REENTER` for TCO, 15+ sites. Frame-based evaluation exists for streaming (`NativeResumeStack`), but typed-apply still recursive.
- **Search:** `SearchContext/ChoicePoint` as seams. `TableStore` for revision-aware memoization. Not full SLG.
- **Storage:** Native and MORK backends via `SpaceMatchBackend`. MORK read-side works; write-side deferred.
- **Formal:** Some Lean infrastructure in mettapedia. Key theorems have `sorrys` or are not started.

This is a mark on the map, not a redrawing of the map.

Part IV

The Horizon

8 Clouds of Possibility

What might we explore next? These are not stages to complete but regions where the landscape is promising and uncertain.

8.1 Eval-Local Term Sharing

Borrowing canonical identity into evaluation arenas could significantly reduce memory pressure on proof-heavy workloads. Positive witnesses include: recursive PLN chaining over medium knowledge bases, metamath proof search at various depths, pattern-heavy stdlib operations, and cross-backend parity tests. No single benchmark defines success; the workloads reveal which optimizations matter.

But it could also introduce lifetime complexity that costs more than it saves. The Conchon–Filliâtre model suggests a path, but the path is not walked.

8.2 Full SLG Tabling

XSB-style suspension/resumption/completion would handle recursive queries that currently loop or produce incomplete answers. The infrastructure is substantial. The question: which workloads need it? Recursive PLN chaining over large knowledge bases is the obvious candidate.

8.3 Dependent Types as Extension

MeTTa’s type system could grow toward dependent types, where types can depend on values. This would enable richer verification and more precise staging. But it also complicates the “is this a value?” question significantly. A cloud to explore, not a near-term commitment.

8.4 `-lang petta` Mode

PeTTa diverges from HE in specific ways. A `-lang petta` flag could enable those divergences intentionally, allowing CeTTa to run PeTTa benchmarks without claiming they’re HE-compliant. Useful for performance comparison and migration.

But `-lang petta` is not the hardest dialect pressure test. Because it stresses reflection, multiset matching, channel/name capabilities, and concurrency semantics, a future `-lang rhocalc` may be the more revealing test of whether the shared search waist is genuinely dialect-parametric rather than merely tolerant of concrete syntax variation.

8.5 PLN/ATP as Optional Lanes

Backward chaining, forward chaining, ATP integration (Lash-style or otherwise)—these are specialized inference modes. They should integrate through explicit lane interfaces, not as modifications to the core evaluator. The `mork:` lane prefix is a start.

8.6 Rho-Calculus Lane

Meredith’s reflective higher-order process calculus [25] provides formal process-algebraic semantics for concurrent computation with reflection. MeTTa’s quote/unquote semantics map naturally to rho-calculus’s `@` and `*` operators, suggesting a principled foundation for concurrent MeTTa execution.

This could enable:

- **Decentralized execution:** RChain-style distributed computation where processes run on different nodes and communicate via channels.
- **Formal concurrency semantics:** No ad-hoc parallelism. The process algebra defines what “concurrent evaluation” means.

- **Reflection as first-class:** In rho-calculus, code-as-data is native, not an encoding. This aligns with MeTTa’s reflective design.

Open question: how much rho-calculus machinery does CeTTa need versus borrow? Full Rholang compilation? A lightweight embedding? A translation layer that preserves key properties?

Reflection also sharpens the future liveness/GC story. If reified spaces can be inspected ambiently, hidden names can leak and naive reachability notions break under lookahead. Any future GC theory for reified spaces should therefore be deployment-relative and capability-disciplined: a process inspects a space only if it was explicitly given that capability. This is a plausible later Lean formalization target once reflective spaces become first-class in the runtime story.

9 The Cloud of Forks

The endgame is not a single design point. It is a region of design space with multiple viable paths. Here we enumerate the major decision axes where multiple options remain open.

9.1 Representation Fork

Option A: Persistent-only hash-consing. Simplest ownership story. Fastest to stabilize. But leaves evaluation-time copying costs.

Option B: Arena-disciplined borrowed sharing. Stronger memory savings. Aligns with one-universe ideal. Needs airtight publication/lifetime rules and formal proof.

Option C: Mostly canonical + local wrappers. Deepest sharing. Strongest long-term symbolic performance. Hardest implementation and proof burden.

9.2 Substitution Fork

Option A: Eager substitution. Simple mental model. Can be expensive on deep substitution chains.

Option B: Selective delayed substitution. Middle ground, focusing on hot cases only.

Option C: Explicit substitutions everywhere. Radically improved copying behavior. Evaluator and normal-form logic become more complex (Abadi et al. $\lambda\sigma$ -calculus).

9.3 Evaluator/Machine Fork

Option A: Direct recursive/trampoline. Clearest semantics, easiest to debug. May suffice for deterministic paths.

Option B: Explicit resumable frame machine. Supports streaming and bounded consumers. Risk of destroying tail-path performance if not careful.

Option C: Derived/compiled machine family. Best performance ceiling. Greatest need for observable-equivalence proofs.

9.4 Tabling Fork

Option A: Revision-aware cache only. Lowest complexity. Insufficient for serious recursion/cycles.

Option B: Variant tabling. Strong step up, simpler than subsumption.

Option C: Subsumptive tabling. More reuse, more subtle semantics/performance tradeoffs.

Option D: Incremental tabling. Strongest fit for changing knowledge bases. Higher complexity.

9.5 Ordering/Demand Fork

Option A: Strict ordered eager results. Simplest user story. Can overcompute badly.

Option B: Ordered but bounded-demand streaming. Likely best near-term target. Pengines-style on-demand.

Option C: Relaxed-order parallel modes. Strongest scalability. Must be explicit semantically.

9.6 Backend/Storage Fork

Option A: Native array/set backend as baseline. Simple, clear benchmark baseline.

Option B: PathMap/trie-heavy backend. Good algebraic/path operations. Risk of conflating membership with identity.

Option C: MORK/graph DB backend. High ceiling for graph-native workloads. Requires strong parity discipline.

Option D: Distributed (DAS) backend. Very large knowledge bases across machines. Explicit cost/consistency model.

Option E: Hybrid multi-backend. Likely true endgame. Requires clean interfaces and parity tests.

9.7 Verification Fork

Option A: Tests only. Insufficient long-term confidence.

Option B: Tests + selected Lean theorems. Practical and high value. The near-term sweet spot.

Option C: Proof-producing kernel / checked delegations. Strongest long-term trust story. More ambitious but strategically valuable.

A Research Annex

Ideas that inspire creativity but aren't on the main map.

Factor-Graph Demand/Supply. Ben’s factor-graph approach to PLN inference, where demand flows backward and supply flows forward through a graph of inference rules. Promising for large-scale reasoning, but not the evaluator core.

ECAN-Style Attention. Economic attention allocation for focusing computational resources on important atoms. Interesting for AGI applications, orthogonal to the runtime’s core concerns.

E-Graphs and Equality Saturation. The `egg` and `egglog` projects show how congruence closure, equality saturation, and Datalog can combine for powerful rewrite optimization. A potential lane, not the evaluator core.

Compiled Datalog (Soufflé). For monotone or stratified logical fragments, Soufflé-style compilation into parallel C++ can yield dramatic speedups. Recognizing when a subproblem has Datalog shape and compiling it aggressively is a long-term optimization opportunity.

OSLF and Type Hypercubes. The OSLF framework [26] shows how to extract a family of type systems (organized as a hypercube) from operational rewrite rules. The modal layer with $\langle K_j \rangle$ modalities and equational center $\mathcal{Z}$ can help classify which operations are pure vs. effectful and which rewrites commute (enabling safe parallelism). MeTTa’s spatial-behavioral type annotations could interface with this framework, deriving effect structure from the rules themselves rather than imposing it externally.

Worst-Case Optimal Joins. For large conjunctive matching, Leapfrog Triejoin [27] and related worst-case optimal join algorithms matter. Variable ordering and multiway join planning become important when MeTTa queries involve multiple patterns over large spaces.

Historical Notes. A git-generated timeline of major commits and version tags would go here. The main roadmap describes territory, not chronology.

B Source Corpus

This section catalogs the primary sources for CeTTa implementation, organized by role. Each entry explains *why it matters* for the endgame.

B.1 Hyperon / MeTTa Semantics

trueagi-io/hyperon-experimental The main public MeTTa implementation. Contains the HE spec in `docs/metta.md`. The semantic reference against which CeTTa must be compared.

Meta-MeTTa (arXiv 2305.17218) Operational semantics for MeTTa. Provides a formal foundation for what evaluation *means*.

Reflective Metagraph Rewriting (arXiv 2112.08272) Conceptual framing for MeTTa as reflective metagraph rewriting. Important for understanding the broader Hyperon vision.

B.2 Comparative Implementations

trueagi-io/PeTTa Efficient Prolog/SWI-based MeTTa. Performance comparator and interoperability target. Shows what WAM-style efficiency looks like.

trueagi-io/metta-wam WAM-targeting MeTTa implementation. Useful for understanding compiled control flow.

logicmoo/metta-testsuite Cross-implementation compatibility harness. Essential for conformance testing.

singnet/das Distributed Atomspace. Large-scale backend with MeTTa-integrated query API. Shows how CeTTa could scale to million-atom KBs.

B.3 Search/Control/Tabling

Warren 1983 The WAM paper. Classic machine model for efficient Prolog. Choicepoints, trail, last-call optimization.

XSB Manual SLG-WAM, tries, call subsumption, incremental maintenance, shared tables. The reference for serious tabling implementation.

SWI-Prolog Tabling Docs Modern production view: JIT indexing, delimited continuations, incremental tabling.

Desouter et al. 2016 “Tabling as a Library with Delimited Control.” Shows how tabling can be implemented cleanly with delimited continuations.

B.4 Term Representation / Indexing

Conchon & Filliâtre 2006 Type-safe modular hash-consing. Canonical reference for safe structural sharing. Weak references, GC integration, lifetime safety.

Graf 1994 Substitution tree indexing. Produces bindings during traversal— exactly what MeTTa pattern matching needs.

Schulz (E Prover) One shared term bank plus multiple specialized indices. The converged architecture after years of “mess.”

Zhu et al. 2025 Modern hash-consing in JuliaSymbolics. 3.2× speedup, 2× memory reduction on real biological models. Validates the payoff.

B.5 Evaluator/Machine Design

Ager et al. (Danvy) “A Functional Correspondence between Evaluators and Abstract Machines.” Derive machine structure from evaluator semantics systematically.

Maurer et al. (“Compiling without continuations”) Join points as first-class optimization concept. Preserves non-escaping tail paths as direct jumps.

Leijen & Lorenzen (“Tail Recursion Modulo Context”) Advanced tail-call optimization that supports structured hot paths without losing abstraction.

B.6 Optional Reasoning Lanes

egraphs-good/egg High-performance equality saturation. E-graph lane for congruence closure and rewrite optimization.

egglog Unifies Datalog and equality saturation. Incremental execution, cooperating analyses, lattice reasoning.

souffle-lang/souffle Highly optimized Datalog with parallel native compilation. Shows how to compile logical fragments aggressively.

Differential Dataflow “Incremental updates with millisecond propagation.” Inspiration for incremental knowledge maintenance.

trueagi-io/hyperon-pln MeTTa-native PLN. The uncertain reasoning lane.

B.7 Verification / Proof

Lean 4 + mathlib Target environment for mechanized theorems. Metaprogramming book for tactic development.

MeTTaMath Bridge between MeTTa and verifiable proof kernels. Metamath-style checking.

vprover/vampllean Proof-producing ATP import into Lean. Model for trusted external reasoning integration.

E Prover, Vampire Mature ATP systems. Potential coprocessors for delegated theorem proving.


TPTP Benchmark/format/evaluation backbone for ATP work. Common problems and standards.

C The Dual-Target Architecture

The mature vision for CeTTa includes a **Lean-style dual-target elaborator**. This is a design direction, not an already proved dual-target implementation. `papers/PureMeTTa.tex` now records the narrower `TwoSortPiSigmaId` experiment and its verified boundaries.

C.1 Proposed Source Syntax $\rightarrow$ Elaborated $\rightarrow$ Proof + Runtime

The intended architecture mirrors Lean’s pipeline: source language would elaborate to a shared typed core with separately justified proof and runtime lowerings. The diagram states that target, not a completed correspondence:



Existing fragment calibration: TwoSortPiSigmaId. The calculus in `Mettapedia/TypeTheory/Calculi/TwoSortPiSigmaId/` has a ground type `u0`, an untyped formation marker `u1`, `Pi/Sigma`, and identity formation/reflexivity. It has no identity eliminator or universe hierarchy. The older permissive typing relation is distinct from the declaration-free `RegularJudgment`, whose subjects are proved strongly normalizing. Exact regular conversion uses computed normal forms; one-pass complete development does not provide full normalization. The MeTTa-specific `Mettapedia/Languages/MeTTa/Experimental/TwoSortPiSigmaId/` adapters establish shared-artifact and restricted operational bridges, not full runtime equivalence. This experiment is not selected as Prime’s proof theory.

CeTTa RuntimeExec (Execution Target). Efficient C runtime with Layer 3 term sharing, hash-consing, PathMap joins, variant tabling. Decidable fragments of the full semantics. This is where performance is achieved.

Elaborated Core (Shared Representation). The proof-of-concept classifier distinguishes four regions: `twoSortKernelRegion` (the experiment’s certificate lane), `runtimeExecRegion` (effectful), `oracleRegion` (FFI/grounded), `metaRegion` (reflection). Sharing this representation does not prove that proof-side reduction and runtime execution agree.

C.2 Semantic Skeleton: CBPV First, MATT Later

The proposed dual-target architecture uses **Call-By-Push-Value** (CBPV, Levy 2001) as a design guide. The following is an intended correspondence, not a proved CBPV model of the two-sort experiment:

CBPV Concept	CeTTa Lane
Value types (things that ARE)	Proposed proof lane: types, proofs, data
Computation types (things that DO)	RuntimeExec: effects, search, space ops
FA (thunk returning A)	Suspend a pure value into runtime
UC (suspended computation)	Embed runtime computation as a value

Key property: Types may depend on *values* but not on arbitrary *computations*. This is exactly what we want: the pure trusted kernel works with values, the runtime produces computations, and the boundary is explicit via the adjunction $F \dashv U$.

Long-term: MATT. CBPV is the 2-mode instance of **Multimodal Adjoint Type Theory** (MATT). As the system matures, we add modes incrementally—Query, Effect, Oracle, Meta—without retrofitting foundations. CBPV is valid MATT from day one.

C.3 Why This Matters

The intended benefits depend on establishing the operational correspondence:

- **Write once, verify and execute:** A future elaborator should support both targets with a proved relation between their results.
- **Proven semantics, optimized execution:** Fragment-specific confluence and typing theorems must be connected to the actual runtime’s sharing and search behavior.

- **Explicit fragment boundaries:** Exact conversion is proved on the regular two-sort normalization domain. That does not decide arbitrary runtime behavior or justify an undecidability claim about this fragment.

C.4 Layer 3 as the Runtime Bridge

CeTTa’s Layer 3 (skeleton + `slot_env`) is precisely the **explicit substitution** representation that bridges elaborated terms to efficient execution:

CeTTa Layer 3	$\lambda\sigma$ -calculus	ρ -calculus
<code>skeleton</code>	M in closure $\langle M, \sigma \rangle$	process with names
<code>slot_env</code>	σ (substitution)	name $\rightarrow$ process map
<code>materialize</code>	$M[\sigma]$	reflection $*x$
<code>canonicalize</code>	closure creation	quotation $@P$
private slot tag	bound variable discipline	ν -restriction

The key theorems (to be proved in `Mettapedia/Bridge/`):

- `materialize_eq_applySubst`: Materialization = substitution application
- `canonicalize_materialize_id`: Roundtrip property
- `cetta_decidable_fragment`: CeTTa is a decidable rho fragment

C.5 Three Sharing Regimes

Layer 3 is *not* the universal term representation. The architecture requires three distinct sharing regimes, each optimized for its context:

Regime	Where	Mechanism
1. Persistent exact	Universe store	Hash-consing, canonical reps
2. Runtime open-term	Evaluator/search	Layer 3 (skeleton + <code>slot_env</code>)
3. Proof-side binder-aware	Proposed proof normalizer	Closures, NbE, substitutions

Anti-goal: Do *not* identify runtime Layer 3 sharing with proof-side normalization. NbE is a possible implementation choice; the existing two-sort regular normalizer is accessibility-driven reduction. A shared artifact alone does not prove equivalence of these representations.

C.6 Seven-Mode Structure

The proposed four-region design could refine into seven modes. This is not an implemented MATTT theorem or a feature list of the fixed two-sort fragment:

1. **Pure** — Trusted kernel: Pi/Sigma/Id/universes, normalization, certificates
2. **Artifact** — Shared typed-core: region classification, overlap certificates
3. **Query** — Read-only: matching, retrieval, bounded search (soundness + completeness)
4. **Exec** — Nondeterministic: rule firing, scheduler, bounded exploration (adequacy)

5. **Effect** — Stateful: add/remove/update space, revision changes
6. **Oracle** — FFI/grounded: explicit contracts, never confused with pure reduction
7. **Meta** — Reflection/tactics: elaboration-time programming, DeLaM-like layers

Each mode has its own typing discipline; modalities connect modes via adjoint functors.

C.7 Four Theorems

Do not ask for one giant bisimulation. Define separate theorems per lane:

1. **ElabSound** — Elaboration yields well-formed artifacts in the appropriate region.
2. **PureAdequacy** — Pure normalization is authoritative: $\Gamma \vdash e : A$ and $e \rightsquigarrow^* v$ implies v is canonical.
3. **RuntimeSound** — Runtime execution is observationally included: $\text{Obs}(\text{exec}(e)) \subseteq \text{Obs}(\text{refSem}(e))$.
4. **RuntimeCompleteOnFragment** — Reverse inclusion on decidable fragments: with bounded parameters, the runtime computes all reachable results.

C.8 Search Parameterization

Full MeTTa evaluation is undecidable. The runtime’s job is to execute *tractable fragments* selected by search parameters:

```
(search-engine depth: 10 breadth: 100 heuristic: bfs backend: pathmap)
```

Parameters split into two classes:

Semantic parameters (affect results):

- **depth:** N — bounds reduction depth
- **breadth:** M — bounds nondeterministic branching
- **fairness:** F — completeness class (DFS, BFS, iterative deepening)
- **fuel:** K — total step budget (opt-in for debugging only)

Pragmatic parameters (affect only efficiency):

- **heuristic:** H — exploration strategy (when extensional results preserved)
- **backend:** B — storage/indexing engine (PathMap, flat, etc.)
- **memo:** M — memoization strategy
- **scheduler:** S — scheduling heuristic

The Pure target specifies *what correctness means*; the Runtime target specifies *what’s computable given resources*. Pragmatic parameters must never change the result set; semantic parameters define the decidable fragment.

C.9 Rholang and Efficiency

A critical observation: **no production rholang implementation has term-level sharing**. RChain’s RSpace hashes *channels*, not *terms*. CeTTa’s Layer 3 architecture could become the first efficient implementation of a rho-calculus fragment with:

- Hash-consing at the term level
- Alpha-canonical variable sharing
- Static resolution of all reflections
- Bounded search with parameterization

This is not incremental improvement—it’s filling a gap that doesn’t currently exist in the rho-calculus implementation landscape.

D HE MeTTa Specification Summary

This section summarizes the Hyperon Experimental (HE) MeTTa specification from `docs/metta.md` in the `trueagi-io/hyperon-experimental` repository. The Lean formalization in `Mettapedia/Languages/MeTTa/` mechanizes these definitions.

D.1 Syntax (S-Expression Grammar)

```
METTA ::= { [ '!' ], ATOM } ;
ATOM ::= SYMBOL | VARIABLE | GROUNDED | EXPRESSION ;
SYMBOL ::= WORD ;
VARIABLE ::= '$', WORD ;
GROUNDED ::= STRING | <tokenizer-matched WORD> ;
EXPRESSION ::= '(' { ATOM } ')'
```

Meta-types: Type, %Undefined%, Atom, Symbol, Variable, Expression, Grounded.

Special forms:

- (= <pattern> <body>) — function definition
- (: <atom> <type>) — type annotation
- (-> $T_1 \dots T_n R$) — function type

D.2 Core Evaluation Functions

metta(\$atom, \$type, \$space, \$bindings) Entry point. Dispatches by meta-type: symbols/grounded go to `type_cast`, expressions to `interpret_expression`.

interpret_expression Checks if head has function type. If applicable, calls `interpret_function` then `metta_call`. Otherwise, treats as tuple via `interpret_tuple`.

interpret_function Evaluates arguments according to declared strictness (from type). Atom type means lazy (unevaluated); other types mean strict.

interpret_tuple Evaluates head, then recursively evaluates tail, assembles result.

metta_call Grounded dispatch or equation matching. Returns equation bodies with substitutions applied.

D.3 Matching Algorithm

Bidirectional atom unification returns a list of binding sets:

match_atoms(\$left, \$right) Recursive structural matching. Symbols must be equal; variables capture values or record equality; expressions match element-wise.

merge_bindings(\$left, \$right) Combines two binding sets, iterating over assignments and equalities from \$right and integrating them into \$left.

add_var_binding(\$bindings, \$var, \$val) If \$var is unbound, assigns \$val. If bound, matches existing value against \$val (recursive).

add_var_equality(\$bindings, \$a, \$b) Records that variables \$a and \$b must unify. If both already have values, matches those values.

Bindings track both assignments ($\$x \leftarrow \text{val}$) and equalities ($\$x = \y). Variable loops are filtered out of results.

D.4 Minimal MeTTa Instructions

The “assembly language” for MeTTa:

- (eval <atom>) — one evaluation step
- (chain <atom> <var> <template>) — interpret then substitute
- (unify <a> <then> <else>) — bidirectional matching
- (function <body>) / (return <atom>) — block control
- (collapse-bind <atom>) / (superpose-bind ...) — nondeterminism
- (metta <atom> <type> <space>) — recursive interpreter call

D.5 Lean Formalization

The Mettapedia formalization covers:

HE/Syntax.lean Atom inductive, meta-types, well-formedness

HE/Eval.lean metta, interpret_expression, etc.

HE/Space.lean Space operations, pattern matching

HE/DeclarativeSpec.lean Big-step relational semantics

HE/CoreFragment.lean Decidable sublanguage for execution

D.6 MM2 Interface

MM2 (Minimal MeTTa 2) is the MORK execution language:

- **Exec rules:** (mm2_exec priority src-pattern sink-action)
 - **Path atoms:** (path from to) for graph traversal
 - **Priorities:** Integer scheduling order
 - **Step function:** (step! &space n) runs n execution cycles
- The Lean formalization in Languages/ProcessCalculi/MORK/ covers:
- **Syntax.lean** — MM2 syntax (ILP, exec rules)
 - **ThreePhaseRefinement.lean** — priority-scheduled execution
 - **Conformance.lean** — MORK-to-Rust correspondence

References

- [1] S. Conchon and J.-C. Filliâtre. Type-safe modular hash-consing. In *ML Workshop*, 2006.
<https://usr.lmf.cnrs.fr/~jcf/publis/hash-consing2.pdf>

- [2] P. Graf. Substitution tree indexing. In *RTA*, 1994. https://pure.mpg.de/rest/items/item_1834191_2/component/file_1857890/content
- [3] D.H.D. Warren. An abstract Prolog instruction set. Technical Note 309, SRI International, 1983. <https://www.sri.com/wp-content/uploads/2021/12/641.pdf>
- [4] The XSB System, Version 4.0, 2022. <https://xsb.sourceforge.net/>
- [5] S. Schulz. E—a brainiac theorem prover (Version 3.x, 2024). *AI Communications*, 2002 (original). <https://github.com/eprover/eprover>
- [6] B. Desouter, M. van Dooren, and T. Schrijvers. Tabling as a library with delimited control. *IJCAI*, 2016. <https://www.ijcai.org/Proceedings/16/Papers/619.pdf>
- [7] TrueAGI. Hyperon Experimental. <https://github.com/trueagi-io/hyperon-experimental>
- [8] TrueAGI. PeTTa: Prolog-based MeTTa. <https://github.com/trueagi-io/PeTTa>
- [9] SingularityNET. Distributed Atomspace (DAS). <https://github.com/singnet/das>
- [10] P. B. Levy. Call-by-push-value. PhD thesis, Queen Mary, University of London, 2001. <https://www.cs.bham.ac.uk/~pbl/papers/thesisqmwphd.pdf>
- [11] M. Abadi, L. Cardelli, P.-L. Curien, and J.-J. Lévy. Explicit substitutions. *Journal of Functional Programming*, 1(4):375–416, 1991
- [12] A. Vandervorst. PathMap: Prefix-compressed structural sharing. <https://github.com/Adam-Vandervorst/PathMap>
- [13] TrueAGI. MORK: High-performance hypergraph kernel. <https://github.com/trueagi-io/MORK>
- [14] egg: E-graphs good. <https://github.com/egraphs-good/egg>
- [15] Soufflé: High-performance Datalog. <https://souffle-lang.github.io/>
- [16] T. J. Green, G. Karvounarakis, and V. Tannen. Provenance semirings. In *PODS*, pages 31–40, 2007. <https://doi.org/10.1145/1265530.1265535>
- [17] I. S. Mumick, H. Pirahesh, and R. Ramakrishnan. The magic of duplicates and aggregates. In *VLDB*, pages 264–277, 1990.
- [18] S. Abiteboul, R. Hull, and V. Vianu. *Foundations of Databases*. Addison-Wesley, 1995. <http://webdam.inria.fr/Alice/>
- [19] Lean 4 theorem prover. <https://github.com/leanprover/lean4>
- [20] B. C. Smith. Reflection and semantics in Lisp. In *POPL*, 1984.
- [21] N. Amin and T. Rump. Collapsing towers of interpreters. In *POPL*, 2018. <https://dl.acm.org/doi/10.1145/3158140>
- [22] U. Berger and H. Schwichtenberg. An inverse of the evaluation functional for typed λ -calculus. In *LICS*, 1991.

- [23] R. E. Korf. Depth-first iterative-deepening: An optimal admissible tree search. *Artificial Intelligence*, 27(1):97–109, 1985.
- [24] D. P. Friedman, W. E. Byrd, and O. Kiselyov. *The Reasoned Schemer*. MIT Press, 2005.
- [25] L. G. Meredith and M. Radestock. A reflective higher-order calculus. In *ENTCS*, 2005.
- [26] M. Stay, L. G. Meredith, and C. Wells. Generating Hypercubes of Type Systems. 2026.
- [27] T. L. Veldhuizen. Leapfrog Triejoin: A Simple, Worst-Case Optimal Join Algorithm. In *ICDT*, 2014. <https://openproceedings.org/2014/conf/icdt/Veldhuizen14.pdf>
- [28] M. Abadi, L. Cardelli, P.-L. Curien, and J.-J. Lévy. Explicit substitutions. *Journal of Functional Programming*, 1(4):375–416, 1991.
- [29] SWI-Prolog Pengines: Web Logic Programming Made Easy. [https://www.swi-prolog.org/pldoc/doc_for?object=section\(%27packages/pengines.html%27\)](https://www.swi-prolog.org/pldoc/doc_for?object=section(%27packages/pengines.html%27))