

Predictive Coding: What Is Exact, What Is Not, and What It Costs

Oruži*

July 2026 — working note

Abstract

We prove what predictive coding (PC) does and does not compute, and test the consequences. Settling any finite linear-Gaussian residual model is exactly conditional Bayesian inference — the unique energy equilibrium *is* the posterior mean, at arbitrary width, depth, and correlated precision. The equilibrium error recursion is exactly backpropagation’s, but the equilibrium weight gradient is not, and we give the counterexample. Exact backprop-equivalent scheduling is characterized completely: a schedule is exact iff the contributions it misses sum to zero, so when contributions cannot cancel it reduces to a structural criterion — every node’s outgoing edges land on a single rank — strictly weaker than the levelizability one expects. Error-based PC is an exact nonlinear change of coordinates from state-based PC, with equal energy and local parameter updates; its first-arrival signal carries exactly the claimed λ^{L-i} factor. Composing that factor with the local update and loss laws proves a depth-credit no-go: inference rates equalize every layer only at the unit-rate backprop boundary, while polynomial Depth- μ P scaling cannot cancel exponential attenuation. Settling on a depth- n path contracts at exactly $\cos(\pi/n)$. For a genuinely joined backbone–adapter path, restricting the coupled operator to a frozen backbone yields exactly the adapter-only first-mode rate; leaving the backbone active re-imports its depth bound. A preconditioned PC step multiplies the loss by $(P-1)^2$: descent holds exactly on $0 < P < 2$, only $P = 1$ reaches zero in one step, and asymptotically all that is required is $\rho(I - P) < 1$ — which a nonnormal Jordan family shows can coexist with arbitrarily large finite transients. The unrestricted conjecture that every saddle of equilibrated PC energy is strict is false: we give a critical cubic-order saddle with zero curvature in every direction. A separately trained centered-sigmoid PC model has exact and robust entropy turnover, while also proving contraction alone insufficient for reachability. Quadratic PC energy cannot orient even a two-node edge from any observational dataset; ordinary clamping is conditioning in both directions. A uniform intervention — clamp a node and delete its own residual term — separates the directions exactly when their covariance is nonzero. On any excited finite two-node intervention design, after both orientations minimize over their own structural coefficient, the true orientation has strictly lower energy if and only if that covariance is nonzero. Finally, two priors with *equal* cross-entropy can have strictly ordered expected search yields: fit does not imply guidance. All statements are machine-checked in Lean 4. Experimentally, PC never overtook backpropagation in any regime we measured: every from-scratch arm failed its registered bar, and the one ignition that reproduced across seeds had to start from a backprop-equivalent anchor.

*Oruži names the AI research collaborators who carried out this work: Anthropic Claude (Opus 4 and Fable 5) and OpenAI Codex (GPT-5.6-Sol) — formalization, implementation, experiments, and drafting. Research direction and oversight: Zar Goertzel.

1 Introduction

Predictive coding networks train by relaxing latent activities against local prediction errors, then applying purely local weight updates. The hope is a biologically plausible — and practically competitive — alternative to backpropagation, and the literature supplies equivalence results in the settling limit together with a growing engineering effort to make deep PC trainable at all [1, 2, 3, 4, 5].

Those equivalence results are stated informally, and they are sloppiest exactly where the hypotheses matter: *which* quantity matches backpropagation, *under what schedule*, *at what depth*, and *at what cost*. This note answers those four questions with theorems, in models small enough that every hypothesis is explicit and every boundary carries a counterexample: linear chains, finite DAGs, scalar paths, finite categorical priors. Each is machine-checked in Lean 4 and axiom-clean, using only `propext`, `Classical.choice` and `Quot.sound`. To keep the statements readable we give them here as ordinary mathematics; Table 1 maps each one to the name you can look up and check.

Two of the answers surprised us. The natural conjecture about which graphs admit exact scheduling is false in an instructive direction (Section 5), and the question every PC-versus-BP comparison reports on — training loss — provably fails to determine the thing those comparisons care about (Section 11).

Section 12 reports the experiments the theory explains. Every arm fixed its acceptance criteria before running; failed arms are reported as failures with their measured failure modes. The companion report *AlienCoding* carries the full empirical account.

2 The model

A predictive-coding network attaches to each node i of a computation graph a latent state z_i , a prediction $f_i(z_{\text{pa}(i)}; w)$ from its parents, and a precision $\lambda_i > 0$. The energy is

$$E(z, w) = \sum_i \lambda_i (z_i - f_i(z_{\text{pa}(i)}; w))^2.$$

Settling clamps the inputs (and, in supervised mode, the outputs) and descends E in z ; the *weight update* is the local gradient $-\partial E/\partial w$ evaluated at the settled state. Backpropagation instead differentiates the supervised loss through the forward activations a . Every equivalence claim in the field is, underneath, a claim about the relationship between the settled state $z^\star$ and the forward state a .

Our running instance is the scalar chain of depth d : latents $z_0, \dots, z_d$, links $i < d$ with gain w_i and precision λ_i , residuals and energy

$$r_i = z_{i+1} - w_i z_i, \quad E(z, w) = \sum_{i < d} \lambda_i r_i^2,$$

clamped at $z_0 = x$ and $z_d = y$. We carry no factor of $\frac{1}{2}$, so the weight gradient is $\partial E/\partial w_i = -2 e_i z_i$ where

$$e_i := \lambda_i r_i$$

is the precision-weighted error at link i . The quantity e_i is the one that will turn out to be exact.

3 Settling is Bayesian inference

Read the chain as a linear-Gaussian generative model: each z_{i+1} is $w_i z_i$ plus noise of precision λ_i . Clamping z_0 and z_d conditions on the two ends.

The right level for this is not the chain. Take any finite *affine residual model*: latents u , residual $r(u) = Au + c$ with A injective, and an arbitrary positive-definite residual precision Λ — so residuals may be correlated, and A need not be square. The energy is $r(u)^\top \Lambda r(u)$.

Theorem 3.1 (Settling computes the posterior mean). *For any such model the posterior precision $Q = A^\top \Lambda A$ is positive definite, and the following coincide: u is an energy equilibrium; u is the unique energy minimizer; and u is the conditional multivariate-Gaussian posterior mean. Any clamped state other than the posterior mean is provably not an equilibrium.*

So “settling performs Bayesian inference” is a theorem, not a slogan: the relaxed latents *are* the posterior mean given the evidence. Scalar chains and vector-valued chains — arbitrary width, arbitrary depth, arbitrary positive-definite (hence correlated) residual precision — are corollaries, obtained by exhibiting the chain’s residual map as such an A and checking injectivity; no chain-specific completion-of-squares argument is repeated. At depth two the scalar case recovers the Kalman update and the two-source Gaussian fusion rule, which is where the PC literature’s Bayesian intuition comes from.

This is also the cleanest way to see the trouble in the next section. The posterior mean given a clamped output is precisely *not* the forward pass — conditioning on the answer moves the latents. That is the point of PC, and it is also why its weight gradient is not backpropagation’s.

4 What matches backpropagation, and what does not

4.1 The error recursion is exact

Theorem 4.1 (Equilibrium errors obey backprop’s recursion). *At any clamped equilibrium of the scalar chain, for every i with $i + 1 < d$,*

$$e_i = w_{i+1} e_{i+1}.$$

This is backpropagation’s backward recursion, exactly, with no limit and no approximation: errors travel from the output end to the input end multiplied by the intervening gains. Weight tying preserves it, which is the credit-assignment certificate for recurrent PC.

The same statement survives the move to the graphs real implementations use. Hash-consed evaluation DAGs share subexpressions, and the design question is whether a shared node gets one latent or one per use.

Theorem 4.2 (Shared-node aggregation). *Let G be a finite DAG with one latent per node, edge gains w_e , node precisions λ_v , and errors $e_v = \lambda_v r_v$. At any equilibrium, every unclamped node v satisfies*

$$e_v = \sum_{e: \text{src}(e)=v} w_e e_{\text{tgt}(e)}.$$

That is reverse-mode accumulation: a shared node’s settled error is the gain-weighted sum of the errors of everything it feeds, which is what differentiation does when it sums gradients over the unfolded tree. The negative fixture is the one that matters to implementors: on a graph where node 0 feeds node 1 through two edges, the correct settled error at node 0 is 2, whereas an

implementation retaining a single parent slot records 1 — half the true value. One latent per shared node, all slots summed, is forced. Our C implementation of PC for the tree network (Section 12.1) is built this way, and its numerical oracles test these shapes.

4.2 The gradient identity is false

The frequently repeated stronger claim is that the equilibrium *weight gradient* equals the backprop gradient. It does not, and linearity does not save it.

Theorem 4.3 (Counterexample). *Take the depth-2 chain with $w_0 = w_1 = 1$, $\lambda_0 = \lambda_1 = 1$, input $x = 1$ and clamped target $y = 2$. The forward state is $a = (1, 1, 1)$ and the clamped equilibrium is $z^* = (1, \frac{3}{2}, 2)$. Then*

$$\left. \frac{\partial E}{\partial w_0} \right|_{z^*} = -1 \quad \text{while} \quad \text{backprop's gradient for } w_0 = -2.$$

The root cause is $z_1^* = \frac{3}{2} \neq 1 = a_1$.

Both gradients use the correct error; they evaluate it against different source activations. Backprop multiplies by the forward activation a_1 ; PC multiplies by the relaxed latent z_1^* , which conditioning has already moved toward the answer. On this fixture the two differ by a factor of two, and nothing about the linear case prevents it — by Theorem 3.1 the displacement $z^* \neq a$ is exactly the Bayesian content of settling.

The consequence for the field is a naming discipline. What PC gets for free is the error recursion. The gradient identity has to be bought, and the price is a schedule.

5 Exact gradients live on a schedule — and only on some graphs

5.1 The schedule

Song et al.’s Z-IL result [2] recovers the exact gradient by timing rather than by settling. Index layers by distance from the output, so that forward consistency reads $w_i a_{i+1} = a_i$, let ρ be the output residual, and let backprop’s errors be $\delta_0 = \rho$, $\delta_{i+1} = w_i \delta_i$, so that backprop’s gradient for w_i is $\delta_i a_{i+1}$.

Theorem 5.1 (Exactness on schedule). *If layer i ’s weights are read at inference step $t = i$ — the step at which the error front first reaches layer i , while its source latent still equals the forward activation a_{i+1} — then the scheduled PC update for w_i equals $\delta_i a_{i+1}$ exactly, for every $i < d$.*

The hypothesis is doing all the work. Read the same parameter at free equilibrium instead and you are back at Theorem 4.3. The identity is a property of *when* you look, not of the fixed point.

5.2 Which graphs admit one

Since exactness depends on reading each parameter exactly when its error arrives, the question becomes: on which graphs does “the first arrival” carry the *whole* error? Give each node a rank, let R be the head rank, and let the contribution travelling along edge e arrive at time $R - \text{rank}(\text{tgt}(e))$. Node v ’s first arrival is $\tau(v) = \min_{e: \text{src}(e)=v} (R - \text{rank}(\text{tgt}(e)))$. Call the schedule *exact at v* when the contributions that have arrived by $\tau(v)$ already sum to the full aggregate.

It pays to separate two things the literature runs together: whether a schedule *numerically* gets the right number, and whether it *structurally* waited for everything. Call a node *admissible* at time

t when every contribution it needs has arrived by t , and let the *missed aggregate* be the signed sum of the contributions that have not.

Theorem 5.2 (Exactness, classified). *For any ranked DAG, any arrival map and any per-node read time, the schedule is exact at every node if and only if every node’s missed aggregate is zero — equivalently, iff every node is either admissible or has its missed contributions exactly cancel. A schedule can therefore be exact while being inadmissible, and that happens precisely when some node’s missed contributions cancel.*

Cancellation is not hypothetical: a node missing contributions $+1$ and -1 is inadmissible and numerically exact at once. So “exact” and “waited for everything” are genuinely different properties, and only the second is structural.

They fuse exactly when cancellation is impossible — which is the regime the PC literature actually works in, since parent contributions there are strictly positive. Under that hypothesis exactness reduces to admissibility, and admissibility is decidable by rank alone. The natural conjecture at that point is unit levelizability: every edge steps one rank, $\text{rank}(\text{tgt}(e)) = \text{rank}(\text{src}(e)) + 1$, so everything arrives in lockstep. That is sufficient. It is not necessary, and the true criterion is weaker and more useful.

Theorem 5.3 (Exact schedules, characterized). *Let G be a finite ranked DAG in which every parent contribution is strictly positive. Then the first-arrival schedule is exact at every node if and only if*

$$\forall v, \forall e_1, e_2 \text{ with } \text{src}(e_1) = \text{src}(e_2) = v : \quad \text{rank}(\text{tgt}(e_1)) = \text{rank}(\text{tgt}(e_2)).$$

That is: exactly when no node’s outgoing edges mix target ranks.

Unit levelizability implies the criterion trivially — all targets sit at $\text{rank}(v) + 1$ — but the converse fails, and we exhibit the witness: a graph with uniform long skips is *not* unit-levelized and *is* exactly schedulable. Skips are not the problem; *ragged* skips are. A residual connection in which a parameter feeds both a one-edge shortcut and a three-edge main path violates the criterion and is proven inexact, which is the mechanism behind Salvatori et al.’s observation that the construction is brittle on non-chain topologies [4].

There is also a value-independent signed classification. Group every missed edge by its target u , and let $g_{v,u}^{\text{miss}}$ be the sum of its gains.

Theorem 5.4 (Structural signed exactness). *A schedule is exact for every parent-error assignment if and only if $g_{v,u}^{\text{miss}} = 0$ for every source v and target u . The same theorem holds for messages in any faithful real module. Thus signed skips may cancel structurally, but only within a target group; with target-injective nonzero gains, or common-sign gains, universal exactness reduces to admissibility.*

The cancellation witness is therefore not a lucky choice of error values: two missed edges to the same target, with gains $+1$ and -1 , are universally exact. A single missed nonzero edge is universally inexact. The scalar and vector-valued fixtures prove both sides.

Read constructively, Theorem 5.3 is an architecture design rule:

Skip connections are compatible with exact backprop-equivalent PC scheduling precisely when they do not mix target ranks. Uniform long skips are fine; rank-mixing residuals are not.

This is the one place where our theory tells an architect what to build rather than what to avoid, and it says the space of exactly-schedulable graphs is strictly larger than the field assumed.

6 Error coordinates: what ePC preserves

Goemaere et al.'s ePC [5] replaces latent states by prediction errors. For arbitrary nonlinear scalar layers f_i , fixed input x , and errors e_i , define recursively

$$z_0 = x, \quad z_{i+1} = f_i(z_i) + e_i.$$

Conversely, any input-clamped state stream gives $e_i = z_{i+1} - f_i(z_i)$.

Theorem 6.1 (Exact state/error reparameterization). *These constructions are mutually inverse. On every finite prefix the ePC energy equals the sPC energy after state reconstruction, and the corresponding local parameter update is identical. At the differential level,*

$$d(E_{\text{sPC}} \circ z)_e = (dE_{\text{sPC}})_{z(e)} \circ dz_e.$$

If dz_e has a continuous-linear right inverse, ePC and sPC have exactly the same critical-point condition.

This is a coordinate equivalence, not a blanket claim that every finite-step algorithm has the same trajectory. The initialization rate still controls how far a signal has travelled and how large it is when it arrives.

Theorem 6.2 (First-arrival attenuation). *For the one-edge-per-step recurrence $s_0 = g$ and $s_{k+1} = \lambda s_k$,*

$$s_{L-i} = \lambda^{L-i} g.$$

If $|\lambda| < 1$, then for every positive tolerance the exact real signal eventually falls below it. It need not equal zero: a machine-checked fixture is nonzero while below tolerance. At $\lambda = 1$, a unit signal never falls below any tolerance at most one. A one-step ePC local update is λ times the backprop product, and is exactly that product at unit rate.

The tolerance statement is deliberately not called floating-point underflow. It is an exact real-number decay theorem plus an explicit external detection policy.

7 The cost of depth

Settling is relaxation, and relaxation on a long clamped chain is slow. This is folklore; here is the exact rate. On a path of n segments with both ends clamped, Jacobi settling is $(Jv)_k = \frac{1}{2}(v_{k-1} + v_{k+1})$ at interior nodes and 0 at the ends.

Theorem 7.1 (Exact relaxation rate). *The mode $s_k = \sin(k\pi/n)$ satisfies $Js = \cos(\pi/n)s$, hence*

$$(J^t s)_k = \cos(\pi/n)^t \sin(k\pi/n) \quad \text{for all } t.$$

The spectral gap obeys $0 \leq 1 - \cos(\pi/n) \leq \pi^2/(2n^2)$, and $\cos(\pi/n) \rightarrow 1$ as $n \rightarrow \infty$.

So the slowest mode decays by exactly $\cos(\pi/n)$ per sweep and the gap closes quadratically in depth: doubling depth quarters the progress per settling step. This is the mathematical spine under the engineering literature devoted to escaping it [5].

Real networks are not uniform, so the natural worry is whether one deep block poisons an otherwise shallow architecture. It does.

Theorem 7.2 (Heterogeneous blocks). *For a finite family of blocks, each with its own depth n_i , coupling eigenvalue c_i and resulting exact mode rate, the family’s modal spectral radius ρ satisfies $\rho < 1$ if and only if every block’s rate is contractive; and if any block i has unit coupling, then*

$$1 - \rho \leq \frac{\pi^2}{2n_i^2}.$$

The bound is the pessimistic one: the family’s spectral gap is capped by the *deepest* unit-coupled block, no matter how shallow everything else is. Stability is a conjunction over blocks; speed is a minimum over blocks.

The corresponding escape is real, but only at an exact operator boundary.

Theorem 7.3 (Frozen-backbone restriction boundary). *Join a backbone path of depth $B > 0$ to an adapter path of depth $A > 1$. The resulting length- $(B + A)$ block Jacobi operator is provably not block diagonal at the attachment. Now freeze by embedding an adapter state into that full operator, applying the full step, and projecting while deleting the clamped attachment row. The restricted operator is exactly the depth- A adapter Jacobi operator, and a scalar r is its first-sine-mode multiplier if and only if*

$$r = \cos(\pi/A).$$

Consequently its first-mode gap is independent of B and obeys $1 - r \leq \pi^2/(2A^2)$. If the backbone instead remains active as a unit-coupled block, then the global modal gap again obeys the backbone bound $1 - \rho \leq \pi^2/(2B^2)$.

The independence is derived from restricting the coupled operator, rather than stored in an adapter definition. It is also sharp about the idealization: a depth-three backbone with a depth-two adapter has frozen adapter rate 0, but modal radius $1/2$ when both unit-coupled paths remain active. Thus shallow adapters escape only when freezing really deletes the backbone coordinates from settling; partial freezing or a shared active relaxation can re-import the backbone depth.

The same first-arrival theorem also prices the imbalance that Depth- μ P was designed to mitigate [7]. Let layer i be $L - i$ links from the output and let every inference link transmit at rate λ . Composing the sealed signal, one-step-update, and loss theorems gives the following result.

Theorem 7.4 (Depth-credit boundary). *The effective per-layer preconditioner is*

$$P_i = \lambda^{L-i}, \quad \ell_i^+ = (1 - \lambda^{L-i})^2 \ell_i.$$

For heterogeneous rates it is the suffix product $P_i = \prod_{j>i} \lambda_j$. Equalizing all P_i forces every used rate to equal one, because the output’s empty product is one. Moreover, for $0 < \lambda < 1$, multiplying by any fixed polynomial in depth still tends to zero; exact cancellation requires the exponential reciprocal $\lambda^{-(L-i)}$.

The Depth- μ P multiplier and its prescribed Adam learning-rate factor cancel each other algebraically, but neither changes the inference rate. Thus they repair a width parameterization, not this exponential credit factor; indeed the reported μ PC experiments omit that Adam rescaling because it was empirically untrainable [7]. This is a scoped result, not a claim that μ PC has no value: it identifies exactly which instability its parameterization cannot remove. The same module also resolves a nearby verbal dispute about whether equilibrated activities “barely move”: in a uniform-precision unit chain, node k moves by exactly $(k/L)(y - x)$. Early fixed nodes move little as depth grows, while nodes near the output need not.

8 What a preconditioner can buy

The standard response to slow or misbehaving settling is to rescale the update — precision scaling, natural-gradient-style metrics, or a line search. In the exactly solvable case, what rescaling buys is a single number, and the answer is unexpectedly sharp.

Take one chain link: source activation s , downstream gain g , weight w , prediction gws , target y , residual $r(w) = gws - y$ and loss $\ell(w) = \frac{1}{2}r(w)^2$. Write $\gamma = gs$ for the effective gain and assume $\gamma \neq 0$. A PC update preconditioned by P moves along the normalized backprop direction, $w^+ = w - P\ell'(w)/\gamma^2$.

Theorem 8.1 (The $(P - 1)^2$ law). $r(w^+) = (1 - P)r(w)$, and therefore

$$\ell(w^+) = (P - 1)^2 \ell(w).$$

Consequently, for $r(w) \neq 0$:

1. $\ell(w^+) = 0$ if and only if $P = 1$;
2. $\ell(w^+) < \ell(w)$ if and only if $0 < P < 2$;
3. the shortfall against the ideal step is exactly $\Delta\ell(1) - \Delta\ell(P) = (P - 1)^2 \ell(w)$;
4. descent degrades strictly monotonically in the squared departure $(P - 1)^2$.

Every arm of our ignition ladder is a point on this parabola. An update in the correct direction at 0.5% magnitude is $P \approx 0.005$: it descends — it is inside $(0, 2)$ — while retaining $(0.995)^2 \approx 99\%$ of the loss per step. That is not a tuning failure, it is the law. And overcorrecting is worse than useless: $P = 3$ lies outside the descent window and *quadruples* the loss, which is what our precision-scaled arm did before it diverged. The window is $0 < P < 2$ and the target is $P = 1$, with quadratic penalty for missing.

The vector case is where the one-step picture must not be over-read. If the residual is an eigenvector of the preconditioner with eigenvalue μ , the one-step law applies with μ in place of P : $\ell \mapsto (\mu - 1)^2 \ell$. It is tempting to conclude that a preconditioner is correct only where its spectrum sits at 1. That reading is wrong, and the asymptotic truth is more permissive.

Theorem 8.2 (Convergence is a spectral-radius condition). Write $D = I - P$ for the preconditioner's departure from the identity. In any complete normed algebra, $\rho(D) < 1$ if and only if $D^n \rightarrow 0$; hence the preconditioned residual iterate converges to zero exactly when $\rho(I - P) < 1$. No diagonalizability is assumed, so nonnormal operators are covered.

So $P = I$ is the *one-step optimum*, not the requirement. Asymptotically an entire open region works, and one may sit far from the identity on every mode and still converge. For a scalar this is the familiar $|1 - P| < 1$ — the window $0 < P < 2$ of Theorem 8.1 — and the spectral form is what survives to matrices.

Finite-time behavior needs its own theorem.

Theorem 8.3 (Geometric envelope). If $\rho(D) < q < 1$, then there is a finite $C \geq 1$ such that

$$\|D^n\| \leq Cq^n \quad \text{and} \quad \|D^n r\| \leq Cq^n \|r\| \quad (n \geq 0).$$

The constant absorbs the finite nonnormal transient; it cannot in general be determined from the spectral radius alone.

The obstruction is exact for the Jordan family

$$D_{\lambda,c} = \begin{pmatrix} \lambda & c \\ 0 & \lambda \end{pmatrix}, \quad D_{\lambda,c}^n = \begin{pmatrix} \lambda^n & nc\lambda^{n-1} \\ 0 & \lambda^n \end{pmatrix}.$$

Theorem 8.4 (Arbitrarily large stable transients). *$D_{\lambda,c}$ has spectrum $\{\lambda\}$, and is nonnormal whenever $c \neq 0$. For $\lambda = 0$ it satisfies $D^2 = 0$, yet a unit residual is amplified on the first step by exactly $|c|$. Consequently, for every proposed bound $B \geq 0$ there is a nonnormal operator with spectral radius zero, exact two-step convergence, and first-step amplification greater than B .*

The earlier witness $c = 2$ is therefore one point in an unbounded family, not a curiosity. Asymptotic stability, geometric rate, transient size, and step-wise descent are four different claims. A non-normal preconditioner may honour the first two while violating the fourth, with the third controlled only by additional operator information.

What no amount of spectral latitude buys is a single global knob.

Theorem 8.5 (One preconditioner cannot serve two modes). *There is a two-mode fixture and a single diagonal preconditioner $\text{diag}(\frac{1}{2}, 3)$ that strictly decreases the loss on one eigenmode and strictly increases it on the other, from equal starting losses.*

This is the no-go for scaling one’s way out: a metric that is merely large, or merely uniform, buys nothing, because its eigenvalues are wrong in both directions at once. The bar is spectral rather than magnitudinal — but by Theorem 8.2 it is a bar of membership, $\rho(I - P) < 1$, not of identity. Information-geometric preconditioning therefore keeps an honest target that is not merely “be a second-order method”.

Line search fares better, because it targets P adaptively rather than fixing it.

Theorem 8.6 (Backtracking descends). *Under an explicit smoothness bound, Armijo backtracking terminates — some $P = P_0\sigma^n$ satisfies the Armijo condition — and every accepted step is non-increasing in energy. This instantiates with the exact smoothness constant of the depth-2 PC energy.*

This is the correctness certificate for the settling procedure used by the surviving experimental arms, whose logged energies were monotone in every accepted run.

9 Not every equilibrated-energy saddle is strict

Innocenti et al. prove strictness for important classes of deep-linear equilibrated PC energy and conjecture that every remaining saddle is strict [6]. A strict saddle is a critical point with a direction of negative second curvature. The unrestricted conjecture has a covariance-degenerate boundary that the partial theorem does not cover.

Take a scalar two-layer model and the two examples with inputs $(1, 1)$ and targets $(1, -1)$. Its exact equilibrated energy at weights (a, b) is

$$F(a, b) = \frac{1 + a^2b^2}{2(1 + b^2)}.$$

This formula is derived in Lean from the rescaled batch residual objective, rather than postulated. Every $(a, 0)$ is critical because the input–target covariance vanishes.

Theorem 9.1 (Strict region and non-strict boundary). *At $(a, 0)$, the output-weight direction has second curvature $a^2 - 1$. Hence every $a^2 < 1$ has a strict-saddle certificate, including the published origin case. But at $(1, 0)$ every directional second curvature is zero, while for every $0 < t < 1$,*

$$F(1 - t, t) < F(1, 0) < F(1 + t, t).$$

Thus $(1, 0)$ is a genuine saddle with explicit descending and ascending approach curves, but it is not strict.

So the universal conjecture is false without a nondegeneracy hypothesis. This does not contradict the published strictification theorem: the counterexample lies outside its covered class. It identifies the missing issue more sharply — a cubic-order saddle can survive exactly where the quadratic curvature vanishes.

10 The coupled system

Settling and learning run together, and their interaction has its own stability region. Write the two-timescale linear system with settling precision λ and learning rate η : the settled error is multiplied by $1 - \lambda$ per step, and the weight by $1 - \eta\lambda$.

Theorem 10.1 (Stability region). *Both multipliers are contractive if and only if*

$$0 < \lambda < 2 \quad \text{and} \quad 0 < \eta\lambda < 2,$$

in which case the coupled iteration converges to zero from every initial state. If $\eta\lambda > 2$, the learning eigenmode diverges in norm.

Note the shape: $0 < \eta\lambda < 2$ is the same window as $0 < P < 2$ in Theorem 8.1, which is not a coincidence — both are the condition for a one-step affine contraction. It is also the formal account of the late-run divergence one of our arms exhibited: the product $\eta\lambda$, not either factor, is what must stay inside the window.

The second dynamical fact concerns what settling does to the *shape* of the output distribution, and it is the one our experiments measured directly.

Theorem 10.2 (Entropy turnover). *For a linear signal of any depth $d > 0$ and any soft target $p \in (0, 1)$ with $p \neq \frac{1}{2}$, the settled output’s entropy $H(s)$ after s settling steps satisfies*

$$H(0) > H_b(p), \quad H(d) = H_b(p), \quad H(d+1) > H_b(p),$$

where H_b is binary entropy. Hence H is neither monotone nor antitone; and it is non-monotone if and only if the target is non-uniform. A gradient-trained comparator provably lands strictly below the target entropy.

Under-settle and the output is diffuse; settle exactly and you hit the target; over-settle and it re-diffuses past it. Settling depth is therefore not a “more is better” knob: there is an interior optimum, and it is at the signal depth. Our tree-network campaign observed this curve — including the turnover — before we proved it.

The turnover is not an artifact of linearity. Let the settling map be any monotone contraction fixing zero, with rate $q < 1$ — a class that includes the saturating nonlinearities PC implementations actually use.

Theorem 10.3 (Nonlinear turnover, sufficient conditions). *Fix such a settling map, a depth $d > 0$, and a non-uniform target $p \in (0, 1)$, $p \neq \frac{1}{2}$. If the depth- d iterate reaches the target logit from the seed, then the same three regimes hold — entropy above target before arrival, equal at depth, above again afterwards — and the entropy is neither monotone nor antitone. Moreover contraction bounds the reachable set: $|\text{logit}(p)| \leq q^d |\text{seed}|$.*

The reachability hypothesis is doing real work and we flag it rather than bury it: this is a theorem about a class of nonlinear settling dynamics that *can* hit their target, not a theorem about arbitrary trained nonlinear PC networks. The contraction bound is the honest sting — it says which targets a given depth can reach at all. We therefore build and train one concrete nonlinear model rather than assume reachability.

Let $\phi(x) = \sigma(x) - \frac{1}{2}$. This centered sigmoid fixes zero, is strictly increasing and globally $\frac{1}{4}$ -Lipschitz, and is proved not to agree pointwise with any real-linear map. Put a parameter w one edge upstream, read $p_w = \sigma(\phi(w))$ after one settling step, choose p_1 as the training target, and train with $L(w) = (p_w - p_1)^2$.

Theorem 10.4 (Trained nonlinear turnover). *$w = 1$ is the unique zero-loss parameter and a global minimizer. At that trained weight, the output probabilities at settling depths 0, 1, 2 are $\frac{1}{2}, p_1, \frac{1}{2}$, with $p_1 > \frac{1}{2}$, so entropy strictly decreases and then increases. The middle entropy remains strictly below uniform throughout an explicit open probability neighborhood of p_1 . The dynamics lift coordinatewise to every finite vector width; an all-one seed gives turnover in every coordinate, while an all-zero seed stays uniform.*

This separates three notions that “contractive nonlinear PC” can obscure. Contraction controls sensitivity, training certifies the chosen target, and the one-edge delay causes turnover. None implies the others: the same centered sigmoid is contractive but its range is $(-\frac{1}{2}, \frac{1}{2})$, so no seed can reach logit 1.

11 Fit does not imply guidance

The last result is not about PC at all. It is about what PC-versus-BP comparisons measure, and it is the reason our empirical verdict is stated in the currency it is.

A search-guidance prior θ over k program tokens is sampled i.i.d. within a budget B ; a *hit* is a draw of the target.

Theorem 11.1 (Yield is target mass). *The expected number of target hits is exactly $B \theta_{\text{target}}$, and is strictly increasing in θ_{target} .*

That is unsurprising, and it is the setup for the separation.

Theorem 11.2 (Equal loss, ordered yield). *Let $\theta^{\text{lo}} = (\frac{3}{4}, \frac{1}{4})$ and $\theta^{\text{hi}} = (\frac{1}{4}, \frac{3}{4})$, with outcome 1 the target. Against the soft target $(\frac{1}{2}, \frac{1}{2})$ these priors have equal cross-entropy, yet for every budget $B > 0$ their expected yields are strictly ordered:*

$$\frac{B}{4} < \frac{3B}{4}.$$

Two models, identical loss, threefold difference in search performance. Cross-entropy is symmetric under swapping mass between outcomes; yield is not, because only one outcome is the target. So training loss — the quantity every PC-versus-BP comparison in the literature reports, including ours — provably does not determine search performance.

What does determine it, in this model, is the mass the prior puts on the target: its *decisiveness*. This is why Section 12.1 reports prior entropy alongside loss, and why the entropies are the numbers that moved with performance.

12 The experiments

12.1 Paired self-learning campaign on the tree network

We reproduced Gauthier’s OEIS program-synthesis pipeline byte-faithfully (Standard ML, HOL4, tree neural network in C) and implanted PC *in place* in the C training code, guarded by numerical oracles (analytic golden vectors to 10^{-9} – 10^{-16} ; the BP path byte-identical when PC is off). The learning rule is the only variable.

Single-generation results establish the mechanism. PC’s yield climbs with settling depth — 63%, 67%, 83%, 93% of BP’s newly solved sequences at 4, 8, 16, 32 settling steps — at roughly linear cost, up to $\sim 10\times$ BP’s training time at 32 steps, and then *turns over* at 64 steps. Throughout, the entropy of the search-guidance prior tracked yield at every point, including the turnover, while training loss did not: PC fit the training objective as well as or better than BP while guiding search worse. Measured prior entropies: BP 0.671; PC $1.479 \rightarrow 0.975$ over steps $4 \rightarrow 32$, re-diffusing to 1.154 at step 64. Theorem 10.2 is the shape of that curve and Theorem 11.2 is why loss failed to see it.

The multi-generation campaign makes the causal comparison real: three paired seed-worlds, one frozen random generation-0 each, both arms grown for five self-training generations under identical budgets. Cumulative solved OEIS sequences:

	gen-1	gen-2	gen-3	gen-4	gen-5
BP, world 1	8,752	11,709	13,224	14,566	15,852
BP, world 2	8,361	11,291	13,250	15,054	16,081
BP, world 3	8,627	11,510	13,267	14,793	<i>in progress</i>
PC, world 1	6,465	6,914	9,067	9,486	10,211
PC, world 2	6,093	8,031	8,359	9,195	9,411
PC, world 3	7,548	7,548	8,820	10,111	<i>in progress</i>
PC/BP ratio	78%	65%	66%	65%	61.4% (w1–w2)

BP compounds metronomically: per-generation deltas agree across worlds within $\pm 1.3\%$ at generation 2, endpoints within $\sim 3\%$. PC ignites but advances in erratic bursts — including one genuine zero-progress generation (world 3, generation 2, byte-identical solution set) — and settles between 58% and 65% of BP’s frontier (world 1 64.4%, world 2 58.5% at generation 5). The final three-world verdict follows the pre-registered analysis when world 3’s generation-5 runs seal.

12.2 The seq2seq ignition ladder

On a faithful reimplement of the alien-coding NMT baseline [10] (14-token program language, 512-unit bidirectional-LSTM encoder–decoder with scaled-Luong attention; the trained reference solves 1,860 new OEIS sequences on seed 1, with 222 target-associated hits against a randomized-association null of 6), we ran a seven-arm overfit-32 ignition ladder for PC variants. The registered composite bar: final loss ≤ 0.05 ; at least 30/32 exact greedy decodes; monotone settling energies; and *distinctness* — the final updates must differ measurably from BP’s, median parameter cosine < 0.95 — so that a pass cannot be secretly backpropagation. The BP anchor passes at loss 0.026, 32/32.

arm	result	loss / exact	measured failure mode
fixed-rate ePC	FAIL	105.4 / 0	settling divergence
backtracking, 16-step	FAIL (time)	—	exceeded registered wall-clock
backtracking, 5-step	FAIL	11.2 / 1	BP-direction updates at 0.5% magnitude
precision-scaled	FAIL	21.4 / 0	late-run divergence
staged BP→PC, 4000	FAIL	0.075 / 32	loss bar only; distinctness passed
staged BP→PC, 4500	FAIL	0.044 / 32	distinctness only; updates BP-like
staged, deepened inference	PASS	0.037 / 32	all bars; median cosine 0.80

Three of these failures are theorems. The 5-step arm is $P \approx 0.005$ in Theorem 8.1: correct direction, 99% of the loss retained per step. The precision-scaled arm left the $0 < \eta\lambda < 2$ window of Theorem 10.1. The fixed-rate arm is the same failure without the line search that Theorem 8.6 certifies.

The joint reading of the two staged arms is the sharpest empirical fact in this note: the staged schedule that starts at provable BP-equivalence and walks toward PC *learns to the bar exactly as far as its updates remain BP-like*, and retains PC-distinctness exactly when it misses the bar. The seventh arm broke the tension: continuing the staged run with deepened inference passed every registered bar at once — loss 0.0366, 32/32 exact decodes, monotone settling, genuinely PC-distinct updates (median parameter cosine 0.80, 13/18 tensors below 0.95).

That pass was seed 0 and judged at one terminal snapshot, so a protocol weakness had to be settled before it could be believed: the terminal cosine statistic turns out to swing *within a single fixed stage* — measured medians 0.929, 0.952, 0.322, 0.787, 0.801 at updates 4501 through 4700 — so no one endpoint is a robust mechanism criterion. The replication was therefore pre-registered on an independent seed and on a prospectively fixed five-checkpoint window instead, with the bars unchanged.

It passes. On seed 1, the complete staged schedule (3000 BP-equivalent one-step anchor updates, then fixed transitions to five-step inference, 4700 updates in total) meets every registered bar: final loss 0.0014, 32/32 exact decodes, all 96 logged inference records monotone, and mechanism distinctness at *all five* checkpoints — median of the checkpoint medians 0.228, against a required cutoff of 0.95. The updates are thus far *more* PC-distinct than seed 0’s terminal 0.80, and distinct throughout the window rather than at a chosen point. Staged ePC ignition reproduces.

Two scope limits are registered with it and remain: the schedule begins from a BP-equivalent anchor, so this is not learning from random initialization; and it is overfit-32 scale, not full-scale NMT performance.

13 Verdict

Dead, on the evidence: predictive coding as a drop-in replacement for backpropagation on synchronous digital hardware, for training program-synthesis models from scratch. At small scale it is dominated at equal epochs and vastly dominated at equal wall-clock; at medium multi-generation scale it holds a stable two-thirds frontier deficit with erratic self-training dynamics; on the recurrent architecture, seven registered arms produced exactly one passing design, and it is staged from a BP-equivalent anchor rather than trained from random initialization.

Surviving, with formal or measured support:

1. *Shallow PC adapters under an exact frozen restriction.* The depth pathology depends only on the adapter after the coupled operator is genuinely restricted (Theorem 7.3), and the empirical verdict independently points at this regime. The theorem does not cover a shared

active normalization or partial freezing: its negative half proves that an active unit-coupled backbone re-imports the backbone depth bound.

2. *BP-train, PC-fine-tune.* The staged homotopy arms were the only ones to reach 32/32 exact decodes; the deepened-inference continuation passed every bar on seed 0, and the pre-registered independent-seed replication passed on seed 1 with distinctness at every checkpoint of a fixed window. This is the one place PC ignited reproducibly under criteria fixed in advance. Buying PC’s properties *after* BP’s learning is the empirically supported direction.
3. *Second-order geometry, with the bar set by Theorem 8.5.* The preconditioning failures are exactly characterized, so the cure is exactly specified: a metric whose departure from the identity is a contraction, $\rho(I - P) < 1$, on the modes that matter — a membership condition, not identity with backprop. Merely large, or merely uniform, is proven not to work.
4. *Rank-uniform skip architectures.* Theorem 5.3 says exact PC scheduling survives skip connections that do not mix target ranks — a strictly larger design space than levelizability, and untested.
5. *Asynchronous and neuromorphic settings.* Untested here; nothing in this note bears on them, except that the decisiveness instrument of Section 11 travels to any setting where a learned prior guides search.

14 Limitations and open problems

The models are exact and, where it matters, small: scalar paths, finite priors, expected rather than at-least-one hit counts. Theorem 5.3 assumes strictly positive parent contributions; Theorem 5.2 now covers the signed case, so what remains open there is not the classification — Theorem 5.4 supplies that — but an architecture-level criterion for useful cancellations under realistic learned gains. Theorem 10.3 still needs reachability in its general form; Theorem 10.4 discharges it for one trained centered-sigmoid model, not for arbitrary coupled multivariate training. Theorem 8.3 gives finite-time nonnormal bounds but its constant is existential; sharp computable constants require more operator structure. Theorem 9.1 refutes universal strictness and classifies one critical line, not every equilibrated-energy critical manifold. The ePC equivalence is mathematical over exact reals; implementation-level rounding behavior remains a separate question. The ignition ladder is overfit-32 scale: its purpose is correctness-gated variant triage, not benchmark performance. The distinctness criterion’s placement at the registered final step is a design choice whose alternatives remain untested. All “dead” verdicts are scoped to the regimes actually measured.

A The causal boundary: clamping is not intervention

A.1 Observation is direction-blind

Take centered variables with positive-definite covariance

$$\Sigma = \begin{pmatrix} v_X & c \\ c & v_Y \end{pmatrix}, \quad \Delta = v_X v_Y - c^2 > 0.$$

The fitted $X \rightarrow Y$ quadratic PC model has coefficient c/v_X and residual variance Δ/v_X ; the reverse model has coefficient c/v_Y and residual variance Δ/v_Y . Their residual operators are provably

distinct. Nevertheless, direct algebra reduces both full energies, pointwise for every real pair (x, y) , to

$$\frac{v_Y x^2 - 2cxy + v_X y^2}{\Delta}.$$

The equality therefore lifts to every finite batch of *arbitrary* observations. There is no Gaussian data hypothesis: “Gaussian” names the quadratic energy. More generally, the formal batch theorem is

$$\sum_s E(z_s) = \text{tr}(Q\hat{\Sigma}), \quad Q = A^\top \Lambda A, \quad \hat{\Sigma} = \sum_s z_s z_s^\top.$$

Changing the data distribution while retaining this energy cannot expose higher-order orientation information. A non-quadratic residual log-density would change the model’s architecture, not merely its environment.

Ordinary PC clamping does not help. Clamp $X = x$ but retain every residual term and minimize over Y . In both orientations the unique equilibrium is $Y = (c/v_X)x$, the ordinary conditional mean. PC clamping is conditioning, and conditioning is direction-blind.

A.2 One uniform severing operation

Pearl’s intervention replaces the intervened structural mechanism rather than conditioning on its output [8, 9]. At the PC energy level this is one orientation-independent operation: clamp a node and delete the residual term owned by that node. Residual rows are indexed by their owning node, so the operation itself never inspects edge direction.

Theorem A.1 (Intervention certificate). *Under $do(X = x)$, the $X \rightarrow Y$ equilibrium remains $(c/v_X)x$, whereas the $Y \rightarrow X$ equilibrium is 0. They differ if and only if $c \neq 0$ and $x \neq 0$. The two interventional energy functions are*

$$E_{X \rightarrow Y}^{do}(x, y) = \frac{(y - (c/v_X)x)^2}{2(\Delta/v_X)}, \quad E_{Y \rightarrow X}^{do}(x, y) = \frac{y^2}{2v_Y},$$

and some intervention separates them if and only if $c \neq 0$.

At the concrete covariance $\begin{pmatrix} 1 & 1 \\ 1 & 2 \end{pmatrix}$ and $x = y = 1$, both half-normalized observational energies equal $1/2$, while the interventional energies are respectively 0 and $1/4$. The negative fixture is just as important: independent equal-variance variables give $y^2/2$ in both orientations, so intervention does not orient an absent edge. This certificate supplies one bit about one edge; it is not a general causal-discovery theorem.

An energy difference at a selected state is not yet a learning theorem, so we also let both orientations fit. On a finite intervention design $(x_s)_s$, take responses from the true forward mechanism, $y_s = (c/v_X)x_s$. A forward candidate fits its own coefficient β and a reverse candidate fits its own coefficient γ ; both are scored by the same severing operation. Their batch objectives reduce to

$$L_{X \rightarrow Y}(\beta) = \frac{(c/v_X - \beta)^2}{2(\Delta/v_X)} \sum_s x_s^2, \quad L_{Y \rightarrow X}(\gamma) = \frac{(c/v_X)^2}{2v_Y} \sum_s x_s^2.$$

The reverse candidates are structurally distinct as γ varies, but their fitted row is precisely the row deleted by $do(X)$; the second objective’s independence from γ is therefore a theorem about the uniform intervention, not a restriction of the competitor.

Theorem A.2 (Fitted two-node causal selection). *For any finite design with $\sum_s x_s^2 > 0$, both orientation-specific objectives attain their minima. The minimized forward energy is strictly lower than the minimized reverse energy if and only if $c \neq 0$.*

Indeed, $\beta = c/v_X$ gives forward minimum zero, while every γ minimizes the reverse objective at the displayed positive value exactly when $c \neq 0$. At independent equal variance both minima are zero, so selection fails. The conclusion is strictly two-node and assumes interventional responses from the forward linear mechanism, positive residual scales, and an excited finite design; it is not a graph-discovery claim.

A.3 A separate continual-learning bridge

Another causal-coding proposal turns on whether learning one thing and then another is the same as learning both. It is not, and the defect is exactly computable in the quadratic case. For two quadratic tasks with curvatures A_1, A_2 , step size η , and gradient g_1 :

$$\theta_{1 \rightarrow 2} - \theta_{\text{add}} = \eta^2 A_2 g_1(\theta), \quad \theta_{1 \rightarrow 2} - \theta_{2 \rightarrow 1} = \eta^2 [A_1, A_2] \theta,$$

the second for centred tasks. So sequential learning departs from additive learning by the second task’s curvature acting on the first task’s gradient, and the order in which you learn matters exactly to the extent that the curvatures fail to commute.

These are two different phenomena, and we separate them rather than let one stand in for the other. A fixture with oblique tasks has a non-zero commutator and is order-dependent. A fixture in which one task is repeated has a *zero* commutator — no order interference is possible — and is still non-additive, because the same cause is counted twice. On the evidence side that double counting is exact: re-applying a Gaussian contribution yields the correct once-each ledger plus precisely one extra copy of that contribution.

We state the scope plainly, and the formalization enforces it: what is proven is linear and quadratic. The general correspondence between predictive coding and causal coding remains explicitly conjectural beyond these instances, and its status now records that a causal use requires an interventional clamp. No theorem promotes the general correspondence.

Where to find these

All statements are machine-checked in Lean 4, in the Mettapedia library under the namespace `Mettapedia.MachineLearning`. Each is axiom-clean, depending only on `propext`, `Classical.choice` and `Quot.sound`. Files below sit in `NeuralNetworks/PredictiveCoding/` unless marked otherwise.

Table 1: Index of formal statements.

	Lean name	File
Thm 3.1	<code>LinearGaussianOperatorModel.equilibrium_iff_eq_posteriorMean</code>	<code>LinearGaussian-Operator</code>
	<code>matrixPCEquilibrium_iff_eq_conditionalPosteriorMean</code>	<code>MatrixGaussianChain</code>
	<code>matrixPC_not_equilibrium_of_ne_posteriorMean</code>	(negative fixture)
	<code>pcConditionalPosteriorMean_equilibrium</code>	<code>LinearGaussianChain</code> (scalar)
Thm 4.1	<code>equilibriumError_satisfies_backpropRecursion</code>	<code>BackpropExactness</code>
	<code>tiedEquilibriumError_satisfies_bpttRecursion</code>	(weight tying)
Thm 4.2	<code>sharedDAGEquilibriumError_satisfies_parentAggregationRecursion</code>	<code>SharedLatentDAG</code>
	<code>dagMultiParent_parentAggregate_counts_both_slots</code>	(positive fixture)
	<code>dagMultiParent_single_slot_is_incorrect_negative_example</code>	(negative fixture)
Thm 4.3	<code>equilibriumGradient_ne_backpropGradient_example</code>	<code>BackpropExactness</code>
	<code>equilibrium_state_ne_forward_state_example</code>	(root cause)

	Lean name	File
Thm 5.1	zilScheduledGainGradient_eq_backpropGradient zilChainResidual_at_schedule_eq_backpropError zilOffSchedule_freeEquilibriumGradient_failure residualSkip_naiveFirstArrivalGradient_failure_example	ZILExactness (front lemma) (negative) (negative)
Thm 5.2	dagScheduleExactAtEveryNode_iff_all_missed_eq_zero dagScheduleExactAtEveryNode_iff_each_admissible_or_cancellation dagSchedule_exact_not_admissible_iff_hasMissedCancellation dagScheduleExactAtEveryNode_iff_admissible_of_noCancellation	DAGSchedule- Exactness (fusion under no cancellation)
Thm 5.4	dagScheduleUniversallyExactAtEveryNode_iff_targetGain_zero dagFirstArrivalScheduleUniversallyExact_iff_targetGain_zero moduleScheduleUniversallyExactAtNode_iff_targetGain_zero moduleSignedCancellation_universallyExact_positive_example moduleSingleMissedEdge_not_universallyExact_negative_example	DAGSchedule- Exactness ModuleSchedule- Exactness (positive fixture) (negative fixture)
Thm 5.3	dagFirstArrivalScheduleExact_iff_uniformOutgoingTargetRanks uniformLongSkipGraph_firstArrival_exact uniformLongSkipGraph_not_unitLevelized	DAGSchedule- Exactness (witness) (witness)
Thm 6.1	spcErrorStream_epcStateStream epcStateStream_spcErrorStream epcPrefixEnergy_eq_spcPrefixEnergy epcCriticalDifferential_iff_spcCriticalDifferential epcLocalParameterUpdate_eq_spcLocalParameterUpdate	ErrorState- Reparameterization ErrorState- Reparameterization
Thm 6.2	spcFirstArrivalSignal_layer_closedForm spcWavefront_eventually_below_tolerance spcWavefront_nonzero_but_below_tolerance_example spcWavefront_unitRate_not_below_small_tolerance epcOneStepParameterUpdate_unitRate_eq_backprop	ErrorState- Reparameterization (positive) (boundary)
Thm 7.4	spcFirstArrivalParameterUpdate_eq_depthCreditBackprop chainLinkLoss_after_uniformDepthCredit_exact inferenceRateEqualization_forces_unitBackpropBoundary muPCDepthScaling_cannot_cancel_contractingDepthCredit uniformPrecision_equilibrium_sub_forward_exact	MuPCDepthCredit
Thm 7.1	pathSlowMode_eigenrelation pathSlowMode_iterate_exact pathRelaxationRate_gap_bound pathRelaxationRate_tendsto_one	DepthPathology
Thm 7.2	HeterogeneousDepthBlockFamily.modalSpectralRadius_lt_one_iff	Heterogeneous- DepthBlocks
Thm 7.3	HeterogeneousDepthBlockFamily.modalSpectralGap_le_depth_bound coupledBackboneAdapterStep_not_blockDiagonal frozenAdapterRestriction_eq_blockPathJacobiStep_clamped frozenAdapterRestriction_eq_blockPathJacobiStep_of_boundary_clamped isFrozenAdapterSlowModeRate_iff frozenAdapterRestriction_gap_bound activeBackboneAdapter_backbone_gap_bound	FrozenAdapter- Restriction

	Lean name	File
Thm 8.1	partialFreezing_reimports_backbone_negative_fixture chainLinkLoss_preconditionedUpdate_exact	Preconditioned- Learning
Thm 8.2	chainLinkResidual_preconditionedUpdate_exact chainLink_zeroLoss_afterUpdate_iff_identity chainLink_preconditionedUpdate_strictDescent_iff chainLink_lossDecrease_gap_from_identity_exact chainLink_lossDescent_strictly_degrades_with_squaredDeparture unitChain_large_preconditioner_increases_loss_negative_example spectralRadius_lt_one_iff_pow_tendsto_zero complexMatrixPreconditioner_fullSpectrum_convergence nonnormalTransientMatrix_fullSpectrum_stable nonnormalTransientMatrix_not_normal	OperatorStability (nonnormal witness)
Thm 8.3	operatorPowNorm_geometricEnvelope	OperatorStability
Thm 8.4	operatorResidualIterate_norm_geometricEnvelope jordanTransientMatrix_pow_exact jordanTransientMatrix_spectrum jordanTransientMatrix_not_normal	OperatorStability
Thm 8.5	spectralRadius_zero_allows_arbitrary_transientAmplification matrixPreconditionedResidual_loss_perMode_exact	(negative boundary) Preconditioned- Learning
Thm 8.6	mixedModePreconditioner_helps_one_hurts_another_fixture armijoBacktracking_terminates armijoCondition_energy_nonincreasing pcEnergyDepth2_armijoBacktracking_terminates_with_descent	BacktrackingDescent
Thm 9.1	orthogonalDataset_outputDirection_curvature	EquilibratedEnergy- Saddles
Thm 10.1	orthogonalDataset_innerStrip_hasStrictSaddleCertificate orthogonalDataset_counterexamplePoint_all_curvatures_zero orthogonalDataset_counterexamplePoint_twoSidedSaddleCurves orthogonalDataset_universalStrictSaddleConjecture_counterexample coupledTwoTimescale_multipliers_stable_iff coupledTwoTimescale_converges_below_threshold coupledTwoTimescale_diverges_above_learning_threshold	(negative boundary) CoupledDynamics
Thm 10.2	arbitraryDepthLinearEntropyTurnover_general arbitraryDepthTurnoverEntropy_nonmonotone_iff	EntropyTurnover
Thm 10.3	monotoneContractiveNonlinearEntropyTurnover_sufficient MonotoneContractiveSettling.iterate_abs_le monotoneContractiveNonlinearEntropyTurnover_uniform_boundary	EntropyTurnover (contraction bound) (boundary)
Thm 10.4	trainedCenteredSigmoidWeight_is_unique_global_minimizer	TrainedNonlinear- Turnover
Appendix	trainedCenteredSigmoidEntropyTurnover approximateTraining_preservesCenteredSigmoidTurnoverMargin centeredSigmoidContraction_does_not_reach_logit_one trainedCenteredSigmoidVectorTurnover centeredSigmoidVectorZeroSeed_uniform PositiveBivariateCovariance.distinct_orientations_same_ observational_batchEnergy PositiveBivariateCovariance.clamp_is_conditioning_and_direction_ blind PositiveBivariateCovariance.interventionalEquilibria_separate_iff	(boundary) (vector lift) (negative fixture) CausalDirection- Identifiability

	Lean name	File
Thm A.1	PositiveBivariateCovariance.observationally_blind_intervention_certificate	
	PositiveBivariateCovariance.TwoNodeQuadraticResidualModel.batchEnergy_eq_trace_quadraticForm_mul_empiricalSecondMoment	
	PositiveBivariateCovariance.separatesDirectionsByInterventionalEnergy_iff	CausalDirection-Identifiability
	PositiveBivariateCovariance.unitNoiseForwardCovariance_observation_tied_intervention_separates	(positive fixture)
	PositiveBivariateCovariance.independentEqualVarianceCovariance_interventionalEnergy_fixture	(negative fixture)
Thm A.2	PositiveBivariateCovariance.forwardOrientationWinsByMinimizedInterventionalEnergy_iff	CausalDirection-Selection
	PositiveBivariateCovariance.isMinimizedForwardInterventionalBatchEnergy_iff	(unique forward fit)
	PositiveBivariateCovariance.reverseCandidates_distinct_but_intervention_deletes_parameter	(anti-vacuity)
	PositiveBivariateCovariance.unitNoiseForwardCovariance_minimizedIntervention_selects_forward	(positive fixture)
	PositiveBivariateCovariance.independentEqualVariance_minimizedIntervention_tied	(negative fixture)
Appendix	sequentialTwoTaskUpdate_sub_additive_exact	ContinualLearning/QuadraticTwoTask
Thm 11.1	centered_sequential_order_defect_eq_commutator	
	linearTwoTask_causalCoding_separation_crown	
	GaussianEvidence.reused_contribution_exact_excess	ContinualLearning/EvidenceLedger
	linearCausalCodingBridge_crown	
	generalCausalCodingCorrespondence_requires_interventionalClamp	(scope marker)
Thm 11.1	iidExpectedTargetHits_eq_budget_mul_targetMass	SearchGuidance/DecisivenessYield
	iidExpectedTargetHits_strictMono_targetMass	
Thm 11.2	equal_crossEntropy_strictly_ordered_expectedYield	SearchGuidance/DecisivenessYield

Related work

Equivalence-in-the-limit results and their scope: [1] (MLPs), [3] (arbitrary graphs, including LSTMs at small scale), [2] (exact scheduling), [4] (topological limits). Engineering of deep PC: [5], whose diagnosis that naive settling misbehaves in digital simulation Theorems 7.1 and 8.1 corroborate from the analytic side. Equilibrated-energy landscape results and the universal strict-saddle conjecture are due to [6]; Section 9 proves a degenerate boundary outside their covered class. The program-synthesis substrate is Gauthier’s OEIS self-learning system and its NMT variant [10]. Our companion report *AlienCoding* documents the replication methodology, the full campaign, and the reference numbers.

References

- [1] J. Whittington and R. Bogacz. An approximation of the error backpropagation algorithm in a predictive coding network with local Hebbian synaptic plasticity. *Neural Computation*, 29(5),

2017.

- [2] Y. Song, T. Lukasiewicz, Z. Xu, and R. Bogacz. Can the brain do backpropagation? Exact implementation of backpropagation in predictive coding networks. *NeurIPS*, 2020.
- [3] B. Millidge, A. Tschantz, and C. Buckley. Predictive coding approximates backprop along arbitrary computation graphs. arXiv:2006.04182, 2020.
- [4] T. Salvatori et al. Learning on arbitrary graph topologies via predictive coding. *NeurIPS*, 2022.
- [5] C. Goemaere, G. Oliviers, R. Bogacz, and T. Demeester. ePC: Fast and deep predictive coding in digital simulation. arXiv:2505.20137, 2025.
- [6] F. Innocenti, E. M. Achour, R. Singh, and C. L. Buckley. Only strict saddles in the energy landscape of predictive coding networks? *Journal of Statistical Mechanics: Theory and Experiment*, 2025. doi:10.1088/1742-5468/ade2eb.
- [7] F. Innocenti, E. M. Achour, and C. L. Buckley. μ PC: Scaling predictive coding to 100+ layer networks. arXiv:2505.13124v2, 2025.
- [8] J. Pearl. The do-calculus revisited. In *Proceedings of the Twenty-Eighth Conference on Uncertainty in Artificial Intelligence*, 2012.
- [9] J. Y. Halpern. Axiomatizing causal reasoning. *Journal of Artificial Intelligence Research*, 12:317–337, 2000.
- [10] T. Gauthier, M. Olšák, and J. Urban. Alien coding. arXiv:2301.11479, 2023.